

Neural Networks and Learning Systems

TBMI26 / 732A55

2024

Lecture 3

Supervised learning – Neural networks

Magnus Borga

magnus.borga@liu.se


Recap - Supervised learning

- **Task:** Learn to predict/classify new data from labelled examples.
- **Input:** Training data examples $\{\mathbf{x}_i, y_i\} \ i=1...K$, where \mathbf{x}_i is a feature vector and y_i is a class label.
- **Output:** A function $f(\mathbf{x}; w_1, \dots, w_N)$ that can predict the class label of \mathbf{x} .


Find a function f and adjust the parameters w_1, \dots, w_N so that new feature vectors are classified correctly. Generalization!!

Linear separability

Linearly separable




Non-linearly separable




Neural Networks

a.k.a. the Multi-layer Perceptron




History of neural networks



- 1960's: Large enthusiasm around the perceptron and "connectionism" (Frank Rosenblatt).
- 1969: Limitations of the perceptron pointed out in a paper by Minsky & Papert, e.g., the XOR problem.
- "Winter period" – little research
- 1980's: Revival of connectionism and neural networks:
 - Multi-layer perceptrons can solve nonlinear problems (this was known before, but not how to train them!)
 - Back-propagation training algorithm
- 1990's: Reduced interest, other methods seemed more promising
- 2010's: Renewed interest – "Deep learning"

A simple 1D example

Training samples with only one feature value!




Not separable with a linear function!


But with a nonlinear function!

$$f(x; w_0, \dots, w_n) = \begin{cases} -1 & |x| > 1 \\ 1 & |x| < 1 \end{cases}$$

Example solution




Nonlinear mapping to a new feature space!



Key: The hidden layer(s)

- The output layer requires linear separability. The purpose of the hidden layers is to make the problem linearly separable!
- **Cover's theorem (1965):** The probability that classes are linearly separable increases when the features are nonlinearly mapped to a higher-dimensional feature space.

The Perceptron Revisited




$$\sigma\left(\sum_{n=1}^N w_n x_n - w_0\right)$$


Minimize the following loss function:

$$\varepsilon(\mathbf{w}) = \sum_{k=1}^{\text{\# training examples}} \left(\sigma\left(\sum_{n=1}^{\text{\# inputs}} w_n x_{nk} - w_0\right) - y_k \right)^2$$

$$\sigma(s) = s$$



$$\sigma(s) = \tanh(s)$$



Some nonlinear activation functions

- Step/sign function

Not differentiable – cannot be optimized!
(by gradient search)

- Hyperbolic tangent

$$\sigma(s) = \tanh(s) \quad \sigma' = 1 - \tanh^2(s) = 1 - \sigma^2$$

- The Fermi-function


$$\sigma(s) = \frac{1}{1 + e^{-s}} \quad \sigma' = \sigma(1 - \sigma)$$

- Gaussian function

$$\sigma(s; \gamma) = e^{-\frac{s^2}{\gamma^2}} \quad \sigma'(s; \gamma) = -\frac{2s}{\gamma} \sigma$$

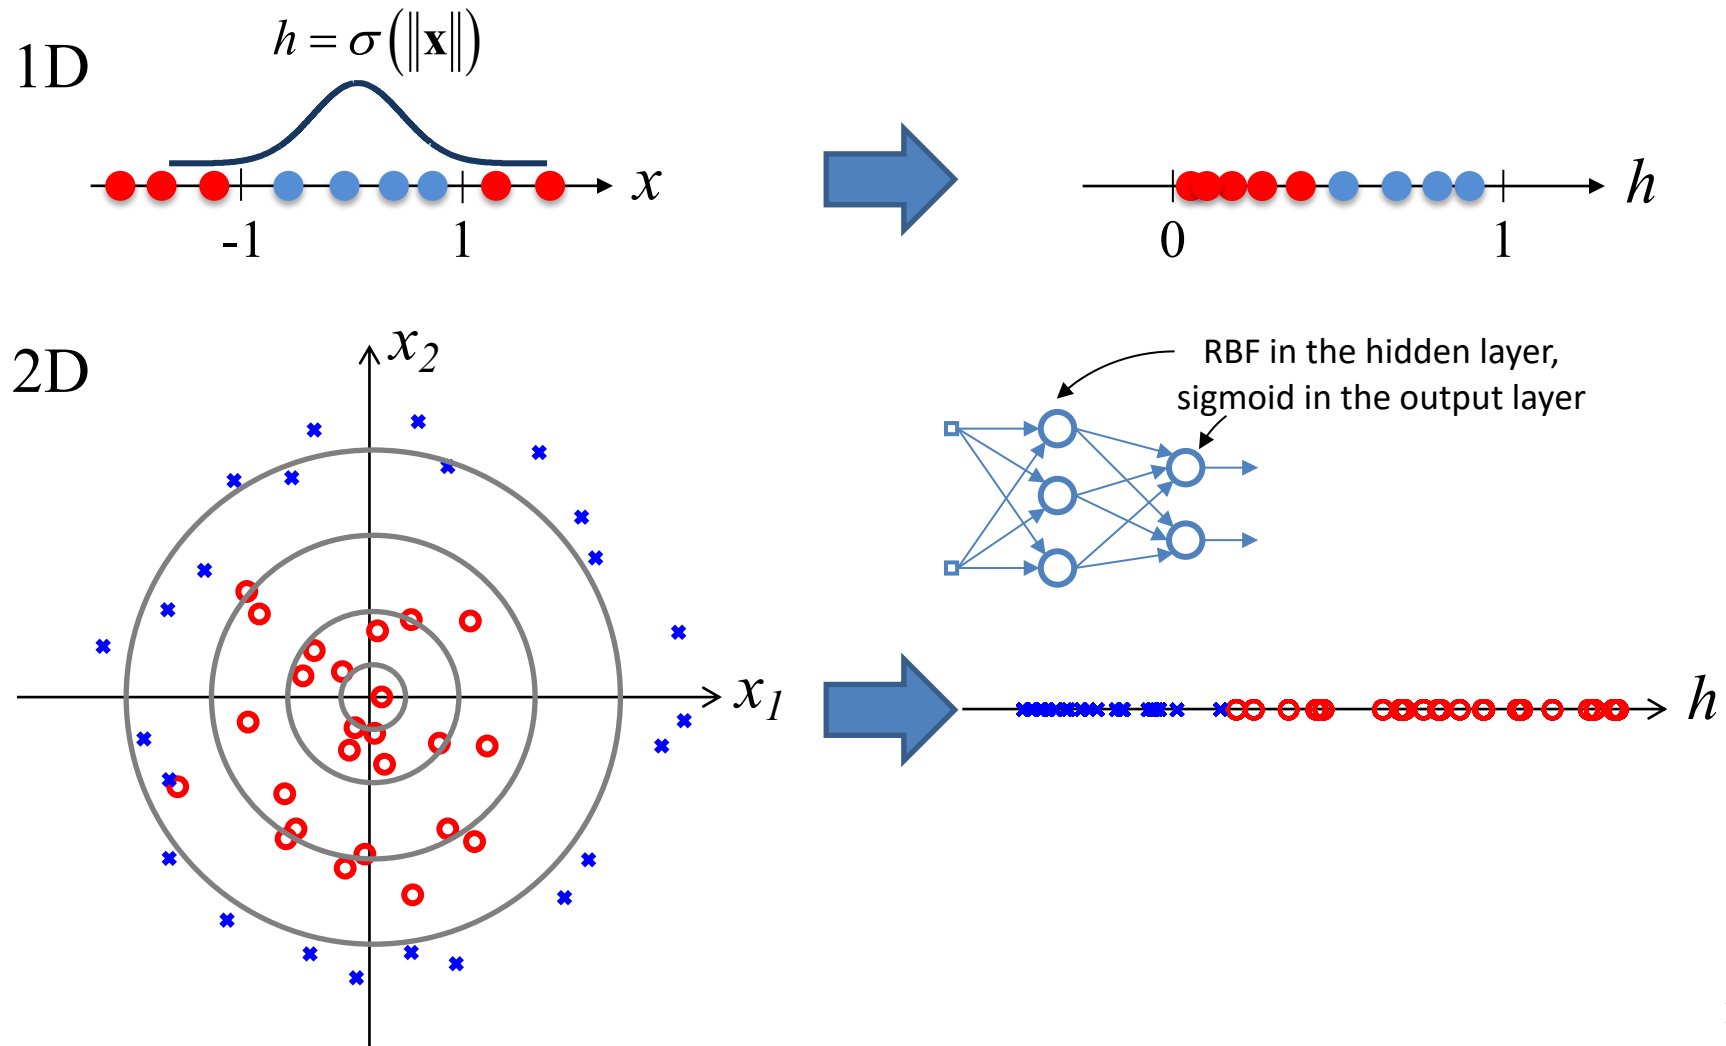
- Rectifying Linear Unit (ReLU)

$$\sigma(s) = \max(0, s) \quad \sigma' = \begin{cases} 0, & s < 0 \\ 1, & s \geq 0 \end{cases}$$





Example - Radial Basis Functions

For example a Gaussian





Example - tanh

$$\sigma(s) = s$$



Same as in
lecture 2!

$$\sigma(s) = \tanh(s)$$



Updated optimization algorithm

$$\varepsilon(\mathbf{w}) = \sum_{k=1}^K \left(\sigma \left(\sum_{n=1}^N w_n x_{nk} - w_0 \right) - y_k \right)^2$$

$$\frac{\partial \varepsilon}{\partial w_i} = 2 \sum_{k=1}^K \left(\sigma \left(\sum_{n=1}^N w_n x_{nk} - w_0 \right) - y_k \right) \sigma' \left(\sum_{n=1}^N w_n x_{nk} - w_0 \right) x_{ik}$$


Gradient descent:

$$\mathbf{w}_{t+1} = \mathbf{w}_t - \eta \frac{\partial \varepsilon}{\partial \mathbf{w}} \quad (\text{Eq. 1})$$

Algorithm:

1. Start with a random \mathbf{w}
2. Iterate Eq. 1 until convergence

Training multi-layer neural networks



Loss function

training examples # output nodes

$$\mathcal{E}(\mathbf{w}) = \sum_{k=1}^K \sum_{m=1}^M \left[y_{mk} - z_{mk}(\mathbf{w}) \right]^2$$

all weights desired output actual output


Stochastic gradient descent

Update using one ($K=1$) training example

$$\mathcal{E}(\mathbf{w}) = \sum_{m=1}^M \left[y_m - z_m(\mathbf{w}) \right]^2$$

$$w_{ij}^{t+1} = w_{ij}^t - \eta \frac{\partial \mathcal{E}}{\partial w_{ij}}$$

From node i to node j
in a layer



The chain rule

$$f(g(x))$$

$$f(g(x), h(x))$$

$$\frac{\partial f}{\partial x} = \frac{\partial f}{\partial g} \frac{\partial g}{\partial x}$$

$$\frac{\partial f}{\partial x} = \frac{\partial f}{\partial g} \frac{\partial g}{\partial x} + \frac{\partial f}{\partial h} \frac{\partial h}{\partial x}$$

Example:

$$f(x; w) = \sigma(\underbrace{wx}_s) \rightarrow \frac{\partial f}{\partial w} = \frac{\partial f}{\partial s} \frac{\partial s}{\partial w} = \sigma'(x, w) \cdot x$$

Note that the activation function σ must be differentiable. This is why the step function does not work with back propagation!

Automatic Differentiation



- Any algorithm that is defined by a sequence of arithmetic operations can be automatically differentiated by repeatedly applying the chain rule!
- This is the basis for error back propagation
- (Note that this is different from numeric and analytical differentiation)

The error back propagation algorithm

- an exercise of the chain rule!

$$\frac{\partial \mathcal{E}}{\partial w_{ij}} = \frac{\partial \mathcal{E}}{\partial z_j} \frac{\partial z_j}{\partial w_{ij}} = \frac{\partial \mathcal{E}}{\partial z_j} \frac{\partial z_j}{\partial s_j} \frac{\partial s_j}{\partial w_{ij}}$$

$$\mathcal{E}(\mathbf{w}) = \sum_{m=1}^M \left[y_m - z_m(\mathbf{w}) \right]^2$$



Back propagation, cont.

$$\varepsilon(\mathbf{w}) = \sum_{m=1}^M \left[y_m - z_m(\mathbf{w}) \right]^2 \quad \frac{\partial \varepsilon}{\partial w_{ij}} = \frac{\partial \varepsilon}{\partial z_j} \frac{\partial z_j}{\partial s_j} \frac{\partial s_j}{\partial w_{ij}}$$


$$\frac{\partial \varepsilon}{\partial z_j} = -2(y_j - z_j)$$

$$\frac{\partial z_j}{\partial s_j} = \sigma'(s_j) = 1 - \sigma(s_j)^2 = 1 - z_j^2 \quad \text{If } \sigma(s) = \tanh(s) \text{ is used!}$$

$$\frac{\partial s_j}{\partial w_{ij}} = h_i \quad (\text{input } i \text{ to unit } j)$$

Updating the hidden layer(s)

$$\frac{\partial \mathcal{E}}{\partial v_{ij}} = ?$$



$$\mathcal{E}(\mathbf{v}) = \sum_{m=1}^M [y_m - z_m(\mathbf{v})]^2$$

A weight in a hidden layer affects **all** output nodes!

$$\mathcal{E}(z_1(\mathbf{v}), \dots, z_M(\mathbf{v}))$$

$$\frac{\partial \mathcal{E}}{\partial v_{ij}} = \sum_{k=1}^M \frac{\partial \mathcal{E}}{\partial z_k} \frac{\partial z_k}{\partial v_{ij}} = \dots \quad \text{Exercise!}$$

Chain rule!

Continue expanding!

Back propagation – Summary


- Two phases:
 - Forward propagation: Propagate a training example through the network
 - Backward propagation: Propagate the error relative to the desired output backwards in the net and update parameter weights.
- Batch update: update after all examples have been presented.

$$\Delta w_{ij} = -\eta \sum_{k=1}^K \frac{\partial \varepsilon(k)}{\partial w_{ij}}$$

Decision boundaries

Neural networks can produce
very complex class boundaries!

$$z(x_1, x_2)$$



$$f(x_1, x_2) = \text{sign}(z(x_1, x_2))$$



What about more than 2 classes?

Multiple outputs each representing likelihood for a certain class



Softmax

$$\sigma(\mathbf{z})_i = \frac{e^{z_i}}{\sum_{j=1}^K e^{z_j}}$$


- Normalized exponential function
- Gives a vector \mathbf{z} where the components' sum is one
- Enables the components of \mathbf{z} to be interpreted as probabilities
- Smooth version of the *argmax* function

Pros and cons of neural networks

- A multi-layer neural network can learn any class boundaries.
- The large number of parameters is a problem:
 - Local optima → suboptimal performance
 - Overfitting → poor generalization
 - Slow convergence → long training times

Overfitting


- The large number of parameters makes it possible to produce overly complicated boundaries.



- A too good fit to the training data can perform poorly for new cases, i.e. worse generalization properties!

Overfitting – Example

Increasing classifier flexibility





Images by B. Schölkopf

Outliers/incorrect labels?


Over-fitting

- The error on training data *always* decreases with increased training
- The error on validation data (the generalization error) decreases in the beginning, but can then start to increase if over-fitting occurs!



Preventing overfitting in neural networks

- **Early stopping:**
Pause training regularly and calculate the performance on the validation data.
- Caution: Validation data becomes training data! Will bias evaluation.
- That's why we need a third dataset for testing – the test data



Training – Validation –Testing

