


Neural Networks and Learning Systems  
TBM126 / 732A55  
2024

**Lecture 7**  
**Reinforcement Learning**

*Magnus Borga*  
*magnus.borga@liu.se*


# Reinforcement learning

Learn by interacting with the environment!



# Examples

Board games




**State:** The board position.

**Action:** Placing a piece in a valid location.

**Reward:** At the end of the game (win/loss).

Robot navigation



**State:** Current location (x,y).

**Action:** Move in a valid direction.

**Reward:** When the way between point A and point B is found. Negative reward for each move that is made.

Chatbots



**State:** Prompt

**Action:** Generated text


**Reward:** Human feedback

# Differences to other methods

- Difference to supervised learning
  - Time!
  - Feedback is given as a scalar reward, not as the correct action to make.
  - Feedback is usually not immediate but is given after many actions – delayed feedback!
  - Can become better than the system designer, unlike a supervised system that can never become better than the teacher.
- Difference to control theory
  - No explicit physical model of the world

# Discretize

Discretize state-space and time!




# Policy

- Defines which action to take in each state
- Can be represented with a look-up table:

State:	$s_1$	$s_2$	$s_3$	$s_4$	....
Action:	$a_4$	$a_1$	$a_{10}$	$a_7$	

# Reward

$r(\mathbf{s}, \mathbf{a})$  – the reward for making action  $\mathbf{a}$  in state  $\mathbf{s}$ .



# Reinforcement learning goal

- The goal in reinforcement learning is to find an optimal policy, i.e., state-action pairs that maximize the reward.
- To do this, we have to solve the following problems:
  1. How to evaluate how good a policy is?
  2. Which policies should we explore?



# Value function

How good is a policy?

A function  $V(s)$  of the state that tells us the value of being in the state given a policy, i.e., the expected amount of reward we get from this state by following the policy:


$$\text{Alt 1: } V(s_t) = \sum_{k=0}^{\infty} r_{t+k}$$

In general, a random variable

$$\text{Alt 2: } V(s_t) = \sum_{k=0}^{\infty} \gamma^k r_{t+k} \text{ where } 0 \leq \gamma \leq 1$$

Makes immediate rewards more important than distant rewards


# Accumulated reward



# Value function, cont.


How good is a policy?

- The value function tells us how good a policy is. The optimal policy has maximum  $V(s)$  for all states.
- When we start learning we explore different policies and we have to learn  $V(s)$  for each policy as it is unknown before we start exploring the environment.
- When using other machine learning approaches, the value function is usually specified by the designer, i.e., not learned by the agent.



# The Credit Assignment Problem

How good is a policy? - How to learn  $V(s)$ ?




## The big question:

How to distribute the reward in the chain of states (and actions) that led to the reward?

For example: Was there a "genius" move that led to the win?

# The Credit Assignment Problem

How good is a policy? - How to learn  $V(s)$ ?




## One solution:

A state is awarded reward according to how many transitions  $m$  from the terminal state it is, i.e., with a discount factor  $\gamma^m$  where  $0 < \gamma < 1$ .

# A Monte Carlo approach

How good is a policy? - How to learn  $V(s)$ ?



- Generate MANY episodes, possibly starting from different states.
- The state-space chains will generally be different because of the stochastic state transitions.
- After a reward  $r$ , update the value functions for all the visited states:

$$\hat{V}(s_k) \leftarrow (1 - \eta) \hat{V}(s_k) + \eta \gamma^m r$$

Learning rate  $0 < \eta < 1$  (points to  $\eta$ )


# states from reward (points to  $m$ )

Discount factor  $0 < \gamma < 1$  (points to  $\gamma$ )

# Temporal Difference approach

How good is a policy? - How to learn  $V(s)$ ?


Update  $V(s)$  after each state transition!



$$\hat{V}(s_k) \leftarrow (1 - \eta) \hat{V}(s_k) + \eta (r_k + \gamma \hat{V}(s_{k+1}))$$

1.  $V(s_k)$  is the expected reward in state  $s_k$ .
2.  $r_k + \gamma V(s_{k+1})$  is a more accurate estimate of the exp. reward in  $s_k$  because we have seen  $r_k$ , and  $V(s_{k+1})$  is closer to the end.
3. Update  $V(s_k)$  with a weighted average using  $\eta$ !

# Reinforcement learning





# How to learn $V(s)$ – Summary

How good is a policy?

- For a given policy, the value (expected reward)  $V(s)$  of each state is unknown before we learn it by interacting with the environment.
- $V(s)$  is found iteratively, starting for example with  $V(s)=0$ , using the Monte Carlo or Temporal Difference methods.
- The Temporal Difference method generally converges much faster.

# But our goal is to find the best policy!

- Brute force – Test all possible policies and choose the one with best value function.
  - Only possible for very small toy problems
- Focus the search more on policies that seem promising, i.e., variations of policies that have already been found to give good value functions.
  - Exploration-Exploitation Dilemma

# Exploration-Exploitation Dilemma

How much should we explore?

- How much should we explore new policies and how much should we exploit what we already learned?
- Example: Multi-armed bandit

Each arm gives a win with a certain probability.

Start by testing both arms to learn which one gives the best win ratio.

When do you stop exploring and start exploiting (pulling only one arm) to maximize your expected winnings?



# Q-function – Add an action dimension


Which policies should we explore?

- Let  $V^*(s)$  denote the value function for the optimal policy.
- $Q(s, a)$ : expected future reward of doing action  $a$  in state  $s$  and then following the optimal policy:

$$Q(s_k, a) = r(s_k, a) + \gamma V^*(s_{k+1})$$

- $Q(s, a)$  indirectly encodes the optimal policy and its value function:  $V^*(s) = \max_a Q(s, a)$ .
- $Q(s, a)$  is unknown – it must be learned!

# Q-learning



# Why the Q-function?

Which policies should we explore?

- Seems that we have just made the problem more complex!?

Added an extra dimension  
in the look-up table!

$Q(s,a)$	$s_1$	$s_2$	$s_3$	...
$a_1$	0.5	0.9	-	...
$a_2$	1.2	-	0.3	...
...	...	...	...	...

The Q-function is an instrument for exploration around the best policies during learning.

# Q-Learning

## Explores policies and value functions simultaneously!

Learn the Q-function using an iteration similar to the Temporal Difference update!

$$\begin{array}{ccc} \text{Updated Q} & \text{Previous estimate} & \text{Better estimate} \\ \hline \underbrace{\hat{Q}(s_k, a_j) \leftarrow (1 - \eta) \hat{Q}(s_k, a_j) + \eta (r + \gamma \hat{V}(s_{k+1}))}_{\text{Updated Q}} = & \underbrace{(1 - \eta) \hat{Q}(s_k, a_j)}_{\text{Previous estimate}} & \underbrace{\eta (r + \gamma \hat{V}(s_{k+1}))}_{\text{Better estimate}} \\ & & \downarrow \\ & & \underbrace{\eta (r + \gamma \max_a \hat{Q}(s_{k+1}, a))}_{\text{Better estimate}} \end{array}$$

# Q-Learning algorithm

Initialize a look-up table for  $Q(s,a)$  with random values.

**for** each episode

    Init a start state  $s$

**repeat** for each step  $k$  in the episode

        Choose an action  $a_j$  (see next slides)

        Take action  $a_j$  and observe reward  $r$  and next state  $s_{k+1}$

        Update estimated  $Q$ -function:

$$\hat{Q}(s_k, a_j) \leftarrow (1 - \eta)\hat{Q}(s_k, a_j) + \eta \left( r + \gamma \max_a \hat{Q}(s_{k+1}, a) \right)$$

**end**

**end**



# $\epsilon$ -greedy exploration

- Make a random action (explore) with probability  $\epsilon$ .
- Make a greedy action (exploit) with probability  $1-\epsilon$ :

$$\arg \max_a \hat{Q}(s, a)$$

- May want to explore more in the beginning (large  $\epsilon$ ) of the training phase and less towards the end.

# Parameter summary

$$\hat{Q}(s_k, a_j) \leftarrow (1 - \eta)\hat{Q}(s_k, a_j) + \eta \left( r + \gamma \max_a \hat{Q}(s_{k+1}, a) \right)$$

- **$0 < \eta < 1$**   
The learning rate. A value close to 0 puts more emphasis on already learned experience and a value close to 1 will overwrite previous experience with new information. Good value to start with: 0.1-0.5
- **$0 < \gamma < 1$**   
The discount factor. A value close to 0 will seek to maximize short-term rewards whereas a value close to 1 will focus the learning on long term rewards. Good value to start with: around 0.9
- **$0 < \epsilon < 1$**   
Exploration factor – the probability of choosing a random action. A value close to 1 will make the learning focus on exploration and a value close to 0 will make the system take actions based on already learned experience. Explore a lot in the beginning (large  $\epsilon$ ) and focus the search more around the good policies towards the end (small  $\epsilon$ ).

# Summary Q-Learning

- Solves the two problems simultaneously:
  1. How to evaluate how good a policy is
  2. Which policies to be explored
- The look-up table quickly becomes (too) big.
- Lots of demos and applets on the internet.

# Reinforcement Learning - Summary

- Learning by trial-and-error
- Often hard to tell *how* a task should be solved but easy to tell *if* and *how well* it has been solved.
- Still mainly a research field
- Strong links to dynamic optimization and optimal control fields.

# A fun example...



[http://www.youtube.com/watch?v=W\\_gxLKSsSIE](http://www.youtube.com/watch?v=W_gxLKSsSIE)