

Neural Networks and Learning Systems  
TBM126 / 732A55  
2024


**Lecture 8**  
**Unsupervised Learning –**  
**Dimensionality Reduction, Clustering and Auto Encoders**

*Magnus Borga*  
*magnus.borga@liu.se*


# Three main categories of machine learning methods

- **Supervised learning (predictive)**  
Learn to generalize and classify new data based on labeled training data.
  - Pattern recognition
  - Classification
- **Reinforcement learning (active)**  
Generate policies/strategies that lead to a (possibly delayed) reward. Learning by doing.
- **Unsupervised learning (descriptive)**  
Discover structure and relationships in complex high-dimensional data.

# Supervised learning – labelled samples



# Unsupervised learning – unlabelled samples



# Unsupervised learning

- **Task:** Find underlying structure in data.
- **Input:** Training data examples  $\{\mathbf{x}_i\} i=1\dots N$ .
- **Output:** Description of the data in a simpler form, e.g., with fewer dimensions or parameters.



# Unsupervised learning

- Optimizes an internal loss function, e.g.
  - max variance (PCA)
  - min distance to cluster centre (k-means)
- Finds a new representation of the data

# Applications

- Clustering
  - find order or structure in data
- Dimensionality reduction / Feature extraction
  - keep the most “important” parts of the signal
    - Image compression
    - image denoising
    - image restoration
    - outlier detection

# Too many dimensions/features


Correlated features or features that do not carry any information:

- Introduce noise in the analysis/classification
- Introduce more parameters in the learning model
  - More local optima in the optimization
  - Poorer generalization
  - Higher computational effort
- Difficult to visualize high-dimensional data

# PCA

## Principal Component Analysis

- Pearson 1901  
(A.k.a. Hotelling-transform or Karhunen-Loéve-transform)
- Coordinate transformation to an orthogonal basis where the data is uncorrelated.
- Dimensionality reduction that preserves maximum variance (minimizes the mean square error).



# Maximize the variance

Normalized vector

The variance in direction  $\hat{\mathbf{w}}$  : (Suppose  $\mathbf{x}$  has mean 0.)

$$\sigma_{\hat{\mathbf{w}}}^2 = E[(\mathbf{x}^T \hat{\mathbf{w}})^2] = E[(\hat{\mathbf{w}}^T \mathbf{x})(\mathbf{x}^T \hat{\mathbf{w}})]$$

$$= \hat{\mathbf{w}}^T E[\mathbf{x}\mathbf{x}^T] \hat{\mathbf{w}} = \hat{\mathbf{w}}^T \mathbf{C} \hat{\mathbf{w}} = \frac{\mathbf{w}^T \mathbf{C} \mathbf{w}}{\mathbf{w}^T \mathbf{w}}$$

The covariance matrix of  $\mathbf{x}$ .

# Maximize the variance

$$\sigma_{\hat{\mathbf{w}}}^2 = \frac{\mathbf{w}^T \mathbf{C} \mathbf{w}}{\mathbf{w}^T \mathbf{w}}$$


$$\frac{\partial \sigma_{\hat{\mathbf{w}}}^2}{\partial \mathbf{w}} = \frac{2}{\mathbf{w}^T \mathbf{w}} (\mathbf{C} \mathbf{w} - \sigma_{\hat{\mathbf{w}}}^2 \mathbf{w}) = 0 \implies$$

$$\mathbf{C} \mathbf{w} = \sigma_{\hat{\mathbf{w}}}^2 \mathbf{w}$$

PCA is the Eigen-value decomposition of  
the data covariance matrix.

# PCA - Example

About 9000 training examples




US Postal Service Digit Data

<http://www.gaussianprocess.org/gpml/data/>

Feature vectors


$$16 \times 16 \rightarrow \mathbf{x} = \begin{pmatrix} x_1 \\ \vdots \\ x_{256} \end{pmatrix}$$

# PCA – Example, cont.







Mean digit

$$\bar{\mathbf{x}} =$$



$$\mathbf{C} = \frac{1}{N} \sum_{i=1}^N (\mathbf{x}_i - \bar{\mathbf{x}})(\mathbf{x}_i - \bar{\mathbf{x}})^T \quad 256 \times 256 \text{ matrix}$$

Eigendecomposition of  $\mathbf{C}$ :



...


256 eigenvectors  
and eigenvalues

50      100      150      200      250


# PCA – Example, cont.

From 256 to 2 feature dimensions!

3 classes




All 10 classes




# Limitations with PCA


Variance is not always the most important goal!



# Class separability



- Small variance
- Large distance



- Large variance
- Small distance

Goal: minimize variance and maximize distance.

# Linear Discriminant Analysis (LDA)

a.k.a. Fishers Linear Discriminant (FLD)

- Minimize variance
- Maximize distance


$$\text{Maximize: } \varepsilon(\mathbf{w}) = \frac{(\mu_1 - \mu_2)^2}{\sigma_1^2 + \sigma_2^2}$$

# LDA – Loss function

$$\varepsilon = \frac{(\mu_1 - \mu_2)^2}{\sigma_1^2 + \sigma_2^2}$$

Distance:

$$\mu(\mathbf{w}) = \frac{1}{n} \sum_{i=1}^n \mathbf{w}^T \mathbf{x}_i = \mathbf{w}^T \left( \frac{1}{n} \sum_{i=1}^n \mathbf{x}_i \right) = \mathbf{w}^T \bar{\mathbf{x}}$$



$$(\mu_1(\mathbf{w}) - \mu_2(\mathbf{w}))^2 = (\mathbf{w}^T \bar{\mathbf{x}}_1 - \mathbf{w}^T \bar{\mathbf{x}}_2)^2 = (\mathbf{w}^T \mathbf{d})^2 = \mathbf{w}^T \underbrace{\mathbf{d} \mathbf{d}^T}_{\mathbf{M}} \mathbf{w} = \mathbf{w}^T \mathbf{M} \mathbf{w}$$

Variance:

$$\sigma^2(\mathbf{w}) = \frac{1}{n} \sum_{i=1}^n \left( \mathbf{w}^T (\mathbf{x}_i - \bar{\mathbf{x}}) \right)^2 = \dots = \mathbf{w}^T \underbrace{\left( \frac{1}{n} \sum_{i=1}^n (\mathbf{x}_i - \bar{\mathbf{x}})(\mathbf{x}_i - \bar{\mathbf{x}})^T \right)}_{\mathbf{C}} \mathbf{w} = \mathbf{w}^T \mathbf{C} \mathbf{w}$$

$$\sigma_1^2(\mathbf{w}) + \sigma_2^2(\mathbf{w}) = \mathbf{w}^T \mathbf{C}_1 \mathbf{w} + \mathbf{w}^T \mathbf{C}_2 \mathbf{w} = \mathbf{w}^T \underbrace{\mathbf{C}_{tot}}_{\mathbf{C}} \mathbf{w}$$

Exercise:  
Complete all the steps!

Note the assumption  $\mathbf{C}_1 \sim \mathbf{C}_2$  !

# LDA – Solution

$$\varepsilon(\mathbf{w}) = \frac{(\mu_1 - \mu_2)^2}{\sigma_1^2 + \sigma_2^2} = \frac{\mathbf{w}^T \mathbf{M} \mathbf{w}}{\mathbf{w}^T \mathbf{C}_{tot} \mathbf{w}}$$

Compare with PCA!

This form is called a Rayleigh quotient,  
which is maximized by the largest eigenvector to the  
generalized eigenvalue problem  $\mathbf{C}_{tot}\mathbf{w} = \lambda \mathbf{M}\mathbf{w}$ !

Simplification:  $\mathbf{M}\mathbf{w} = \mathbf{d}\mathbf{d}^T \mathbf{w} = k\mathbf{d}$


$\underbrace{\phantom{...}}_{\text{Some scalar } k}$

$$\mathbf{w} \sim \mathbf{C}_{tot}^{-1} \mathbf{d}$$

Scaling of  $\mathbf{w}$  not important!

# LDA - Summary


- Projection direction  $\mathbf{w}$  on which two classes are maximally separated



- Closed form solution, no parameters to set, no local optima.
- Components of  $\mathbf{w}$  give the importance of each feature in  $\mathbf{x}$ .
- Assumes identical distributions of the two classes!
- Note that this is actually supervised learning since labelled data is required!



# Categorization

- Categorization and grouping of objects based on similar properties is an important functionality in learning and knowledge representation.
- In the machine learning area, this is usually referred to as *clustering*.



# What describes a cluster?

- Distances to other points?
- Connectivity?
- Different definitions lead to different algorithms.



# $k$ -Means algorithm

- Assume  $k$  clusters (user defined).
- Represent each cluster with a mean prototype vector  $\mathbf{p}_j$  at the cluster center.
- A data point belongs to the cluster with the closest prototype vector (Euclidian distance).






# $k$ -Means algorithm, cont.

1. Start with  $k$  random prototype vectors  $\mathbf{p}_j$
2. Iterate:
  1. Assignment: Assign each data vector  $\mathbf{x}_i$  to the closest prototype vector  $\mathbf{p}_j$ . Denote the set of data vectors assigned to cluster  $\mathbf{p}_j$  by  $S_j$ .
  2. Update all prototype vectors to the mean of the clusters:

$$\mathbf{p}_j = \frac{1}{|S_j|} \sum_{k \in S_j} \mathbf{x}_k$$


Number of elements in  $S$

# $k$ -Means - Example




# $k$ -Means - Example


Random initialization



1st iteration




2nd iteration




T1 MRI



T2 MRI



$k$ -means result



# $k$ -Means - Discussion

- Must specify  $k$
- Tries to minimize the loss function

$$\varepsilon(\mathbf{p}_1, \dots, \mathbf{p}_k, S_1, \dots, S_k) = \sum_{i=1}^k \sum_{\mathbf{x}_j \in S_i} \|\mathbf{x}_j - \mathbf{p}_i\|^2$$

- Note that this function is not differentiable as the  $S_i$ :s are discrete sets, i.e., we cannot do gradient descent.

# Expectation Maximization (EM) – Intro

$$\varepsilon(\mathbf{p}_1, \dots, \mathbf{p}_k, \underbrace{S_1, \dots, S_k}_{\text{Unknown class labels}}) = \sum_{i=1}^k \sum_{\mathbf{x}_j \in S_i} \|\mathbf{x}_j - \mathbf{p}_i\|^2$$

Assume we know  $S_1, \dots, S_k$ !

$$\frac{\partial \varepsilon}{\partial \mathbf{p}_i} = \frac{\partial}{\partial \mathbf{p}_i} \left( \sum_{\mathbf{x}_j \in S_i} \|\mathbf{x}_j - \mathbf{p}_i\|^2 \right) = -2 \sum_{\mathbf{x}_j \in S_i} (\mathbf{x}_j - \mathbf{p}_i) = 2|S_i| \mathbf{p}_i - 2 \sum_{\mathbf{x}_j \in S_i} \mathbf{x}_j$$

$$\frac{\partial \varepsilon}{\partial \mathbf{p}_i} = 0 \quad \xrightarrow{\text{Mean vector!}} \quad \mathbf{p}_i = \frac{1}{|S_i|} \sum_{\mathbf{x}_j \in S_i} \mathbf{x}_j$$

# Expectation Maximization (EM) – Intro

$$\varepsilon \left( \underbrace{\mathbf{p}_1, \dots, \mathbf{p}_k}_{\text{Continuous parameters}}, S_1, \dots, S_k \right) = \sum_{i=1}^k \sum_{\mathbf{x}_j \in S_i} \|\mathbf{x}_j - \mathbf{p}_i\|^2$$

Now, assume we know  $\mathbf{p}_1, \dots, \mathbf{p}_k$ !

Each  $\mathbf{x}_j$  independently contributes a distance  $\|\mathbf{x}_j - \mathbf{p}_i\|$  to  $\varepsilon$

Obvious that  $\varepsilon$  is minimized if each  $\mathbf{x}_j$  is assigned to the set  $S$  associated with the closest  $\mathbf{p}$ !

k-Means algorithm!!


# Expectation Maximization

- The  $k$ -means algorithm is a special case of a general optimization approach called *Expectation Maximization* (EM).
- EM can be used when we want to estimate model parameters (the prototypes  $\mathbf{p}_i$ ), but for each data sample  $\mathbf{x}_j$  there is a hidden/missing discrete parameter (the class labels  $S_j$ ).
- EM iterates between optimizing the hidden parameters and the model parameters.

# Mixture of Gaussians (MoG) clustering


Another application of EM!

*k*-means



May not be correctly clustered!

**Idea:** Represent each cluster using a Gaussian distribution



Problem: Each data sample  $\mathbf{x}_j$  belongs to one of  $k$  Gaussian distributions  $N(\mathbf{p}_i, \mathbf{C}_i)$ ,  $i=1..k$ .

Find the sets  $S_i$  of samples that belong to distribution  $N(\mathbf{p}_i, \mathbf{C}_i)$  and the mean  $\mathbf{p}_i$  and covariance matrix  $\mathbf{C}_i$  of that distribution.

# Let's use EM!

Assume we know the hidden parameters  $S_1, \dots, S_k$ !

That is, we know the samples for each set, e.g.,  $S_1 = \{\mathbf{x}_1, \mathbf{x}_7, \mathbf{x}_{12}, \mathbf{x}_{13}\}$

We can then estimate the mean  $\mathbf{p}_i$  and covariance  $\mathbf{C}_i$  using standard estimation for the Gaussian:

$$\mathbf{p}_i = \frac{1}{|S_i|} \sum_{k \in S_i} \mathbf{x}_k$$

$$\mathbf{C}_i = \frac{1}{|S_i|} \sum_{k \in S_i} (\mathbf{x}_k - \mathbf{p}_i)(\mathbf{x}_k - \mathbf{p}_i)^T$$

# Let's use EM, cont!

Assume now that we know the Gaussian distribution parameters, i.e., we know all distributions

$$f(\mathbf{x}; \mathbf{p}_i, \mathbf{C}_i) = \frac{1}{(2\pi)^{d/2} |\mathbf{C}_i|^{1/2}} e^{-\frac{1}{2}(\mathbf{x}-\mathbf{p}_i)^T \mathbf{C}_i^{-1} (\mathbf{x}-\mathbf{p}_i)}$$

Let the hidden parameter  $S_i$  be the set of all data samples for which  $f(\mathbf{x}; \mathbf{p}_i, \mathbf{C}_i)$  is larger than for all other distributions. That is, each data sample is assigned to the Gaussian distribution **to which it most likely belongs!**

# Mixture of Gaussians – Using EM

1. Start with  $k$  random Gaussians  $(\mathbf{p}_j, \mathbf{C}_j)$
2. Iterate:
  1. Assign each sample  $\mathbf{x}_i$  to the most likely Gaussian

$$\max f(\mathbf{x}_i; \mathbf{p}_j, \mathbf{C}_j) = \frac{1}{(2\pi)^{d/2} |\mathbf{C}_j|^{1/2}} e^{-\frac{1}{2} (\mathbf{x}_i - \mathbf{p}_j)^T \mathbf{C}_j^{-1} (\mathbf{x}_i - \mathbf{p}_j)}$$


Denote the set of samples assigned to Gaussian  $j$  by  $S_j$ .

2. Update all Gaussians' means and covariances:

$$\mathbf{p}_j = \frac{1}{|S_j|} \sum_{k \in S_j} \mathbf{x}_k$$

$$\mathbf{C}_j = \frac{1}{|S_j|} \sum_{k \in S_j} (\mathbf{x}_k - \mathbf{p}_j)(\mathbf{x}_k - \mathbf{p}_j)^T$$


# Mixture of Gaussians - Example



# Summary of $k$ -Means and MoG clustering

- Must choose the number of clusters  $k$ .
- Different initializations may give different results.
- May converge to degenerate solutions, e.g., empty clusters.
- MoG allows for elliptic cluster shapes with different sizes.

# Auto Encoders




## Image segmentation

- Training data  $\{x_i, y_i\}$
- Binary output (foreground / background)
- Classification of each pixel

## Auto encoder


- Training data  $\{x_i, x_i\}$
- Does not require labelled training data – Unsupervised learning!
- The smallest layer must be able to represent all training data
- The bottle neck forces an efficient coding of the original data (images)

# Auto Encoders



- The hidden representation  $h$  is a compact encoding of the distribution of training data  $x$
- If the dimensionality of  $h$  is chosen properly,  $h$  can represent all training data well, but not data points that deviate from the distribution of  $x$

# Data manifold



- The data lies in a low-dimensional manifold in a high-dimensional feature space.
- An auto encoder learns to represent the manifold
- Data points that are outside the manifold cannot be accurately represented

# Outlier detection




- Data points that deviate from the distribution of normal samples (the learned manifold)
- These data points can indicate anomalies such as production errors or pathologies

# Application in Medical Image Analysis

- Train on “normal” images
- Compare input and output
- Use the residual as indication of abnormality



Normal  
brain



Brain with  
lesion



Residual  
overlay