

Neural Networks and Learning Systems
TBM126 / 732A55
2023

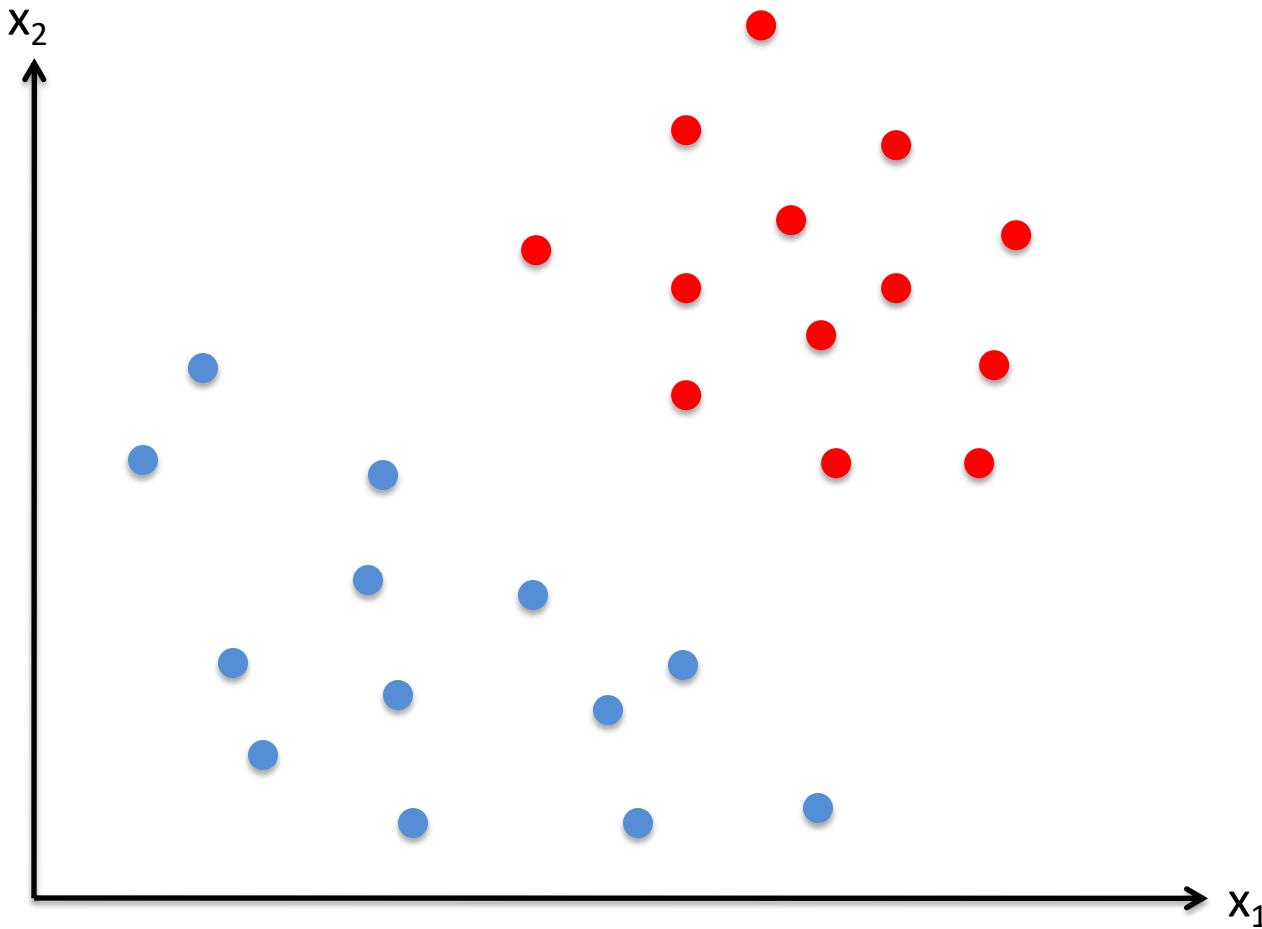
Lecture 8
Unsupervised Learning –
Dimensionality Reduction, Clustering and Auto Encoders

Magnus Borga
magnus.borga@liu.se

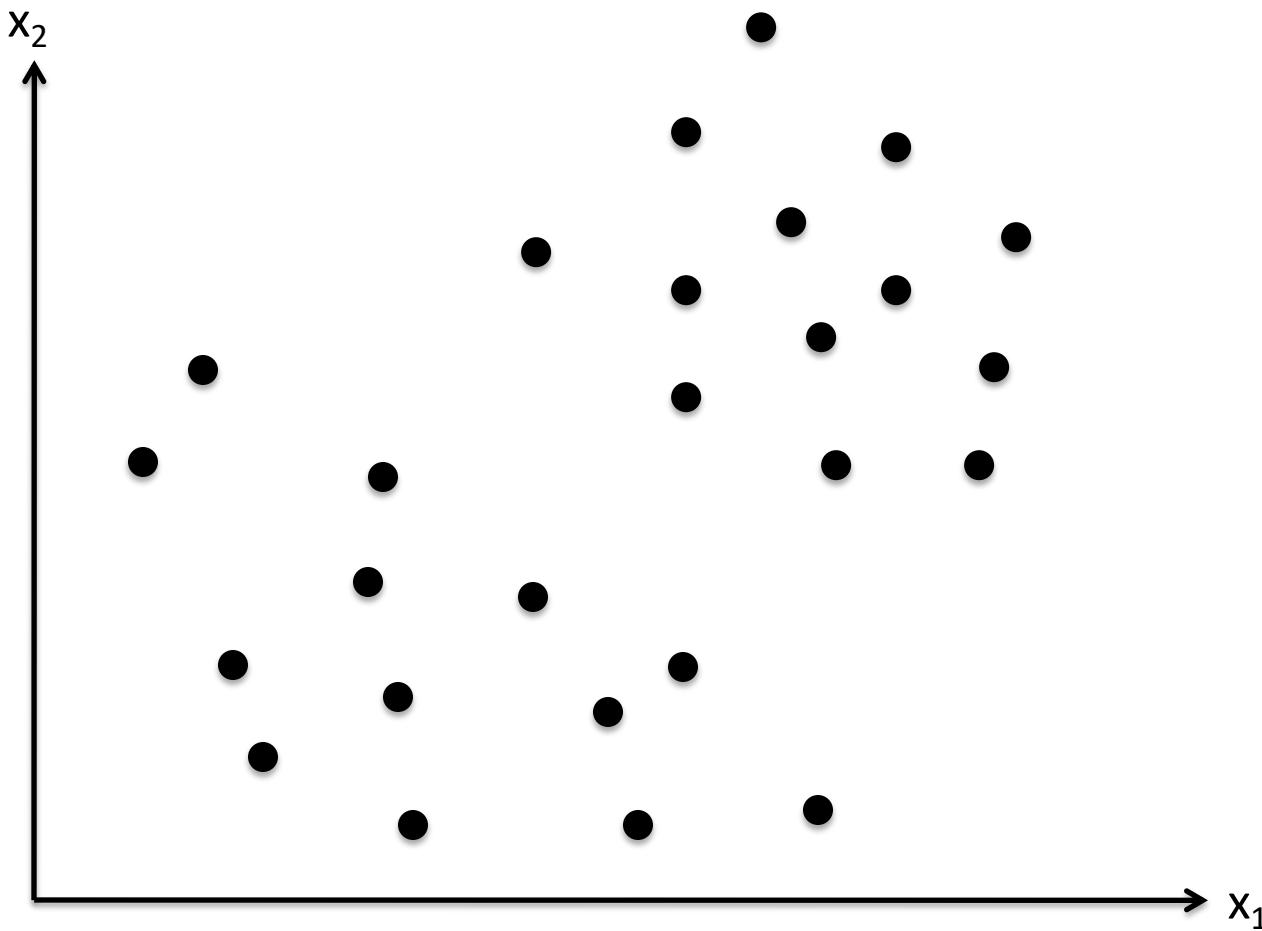
Three main categories of machine learning methods

- **Supervised learning (predictive)**
Learn to generalize and classify new data based on labeled training data.
 - Pattern recognition
 - Classification
- **Reinforcement learning (active)**
Generate policies/strategies that lead to a (possibly delayed) reward. Learning by doing.
- **Unsupervised learning (descriptive)**
Discover structure and relationships in complex high-dimensional data.

Supervised learning – labelled samples

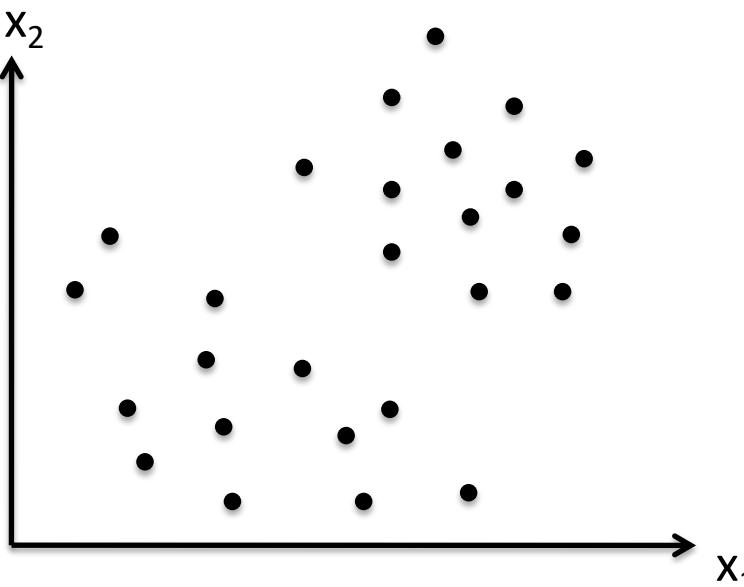


Unsupervised learning – unlabelled samples



Unsupervised learning

- **Task:** Find underlying structure in data.
- **Input:** Training data examples $\{\mathbf{x}_i\}$ $i=1\dots N$.
- **Output:** Description of the data in a simpler form, e.g., with fewer dimensions or parameters.



Unsupervised learning

- Optimizes an internal loss function, e.g.
 - max variance (PCA)
 - min distance to cluster centre (k-means)
- Finds a new representation of the data

Applications

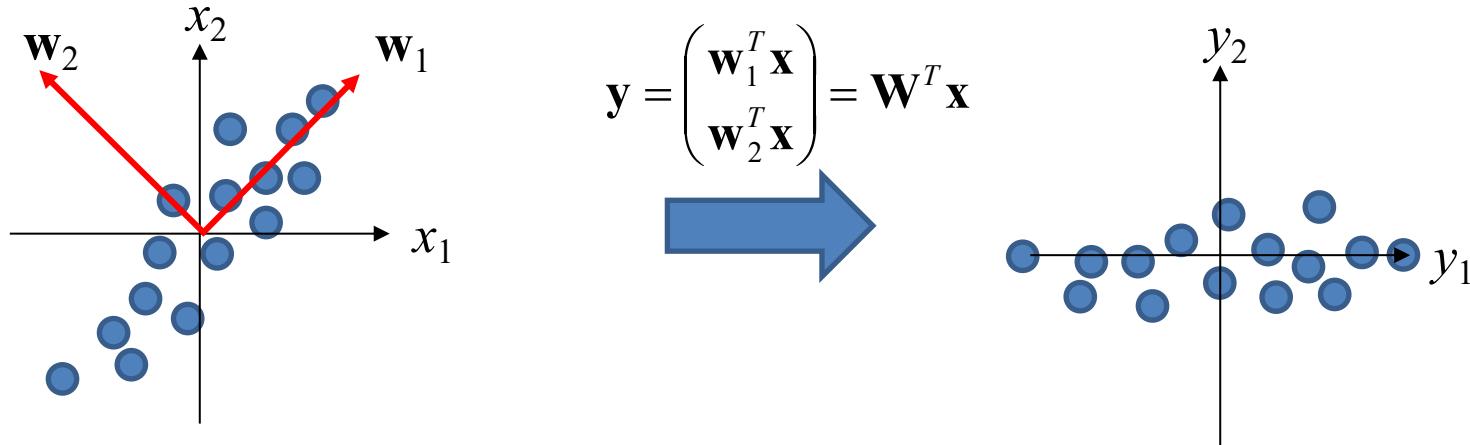
- Clustering
 - find order or structure in data
- Dimensionality reduction / Feature extraction
 - keep the most “important” parts of the signal
 - Image compression
 - image denoising,
 - image restoration,
 - outlier detection

Too many dimensions/features

Correlated features or features that do not carry any information:

- Introduce noise in the analysis/classification
- Introduce more parameters in the learning model
 - More local optima in the optimization
 - Poorer generalization
 - Higher computational effort
- Difficult to visualize high-dimensional data

Projection / linear transformation



Can "uncorrelate" data through a linear transformation!!

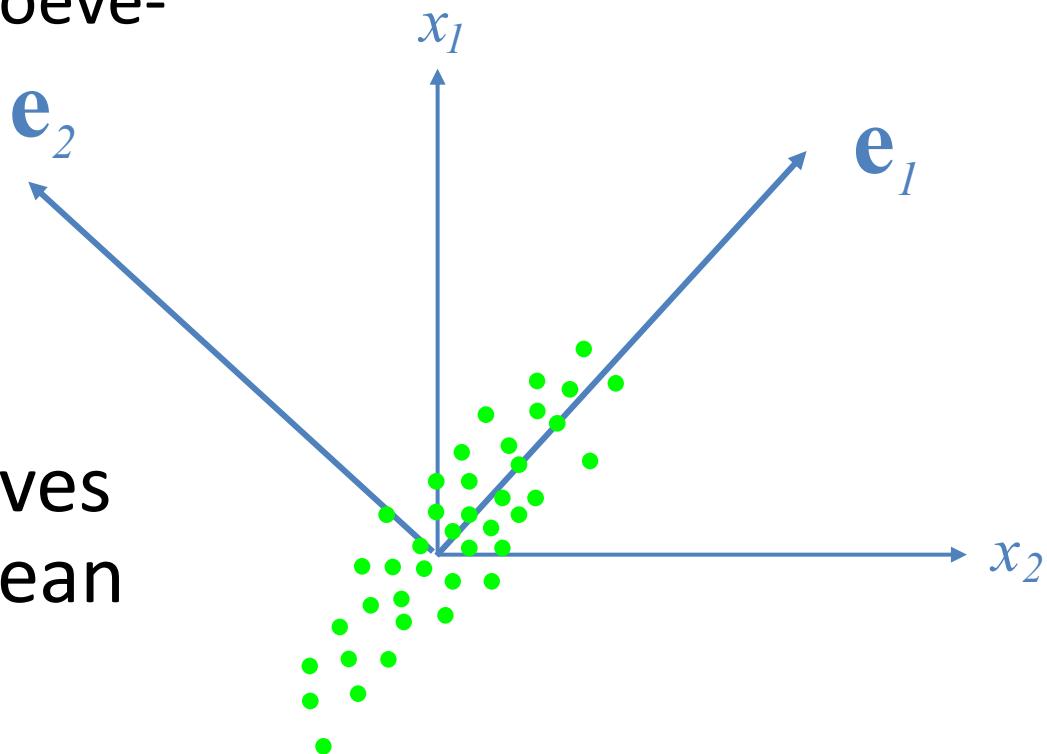
$$\mathbf{C}_x = \begin{bmatrix} \text{Var}(x_1) & \text{Cov}(x_1, x_2) \\ \text{Cov}(x_2, x_1) & \text{Var}(x_2) \end{bmatrix}$$

$$\mathbf{C}_y = \begin{bmatrix} \text{Var}(y_1) & 0 \\ 0 & \text{Var}(y_2) \end{bmatrix}$$

PCA

Principal Component Analysis

- Pearson 1901
(A.k.a. Hotelling-transform or Karhunen-Loéve-transform)
- Coordinate transformation to an orthogonal basis where the data is uncorrelated.
- Dimensionality reduction that preserves maximum variance (minimizes the mean square error).



Maximize the variance

Normalized vector

The variance in direction $\hat{\mathbf{w}}$: (Suppose \mathbf{x} has mean 0.)

$$\sigma_{\hat{\mathbf{w}}}^2 = E[(\mathbf{x}^T \hat{\mathbf{w}})^2] = E[(\hat{\mathbf{w}}^T \mathbf{x})(\mathbf{x}^T \hat{\mathbf{w}})]$$

$$= \hat{\mathbf{w}}^T E[\mathbf{x}\mathbf{x}^T] \hat{\mathbf{w}} = \hat{\mathbf{w}}^T \mathbf{C} \hat{\mathbf{w}} = \frac{\mathbf{w}^T \mathbf{C} \mathbf{w}}{\mathbf{w}^T \mathbf{w}}$$

The covariance matrix of \mathbf{x} .

Maximize the variance

$$\sigma_{\hat{\mathbf{w}}}^2 = \frac{\mathbf{w}^T \mathbf{C} \mathbf{w}}{\mathbf{w}^T \mathbf{w}}$$

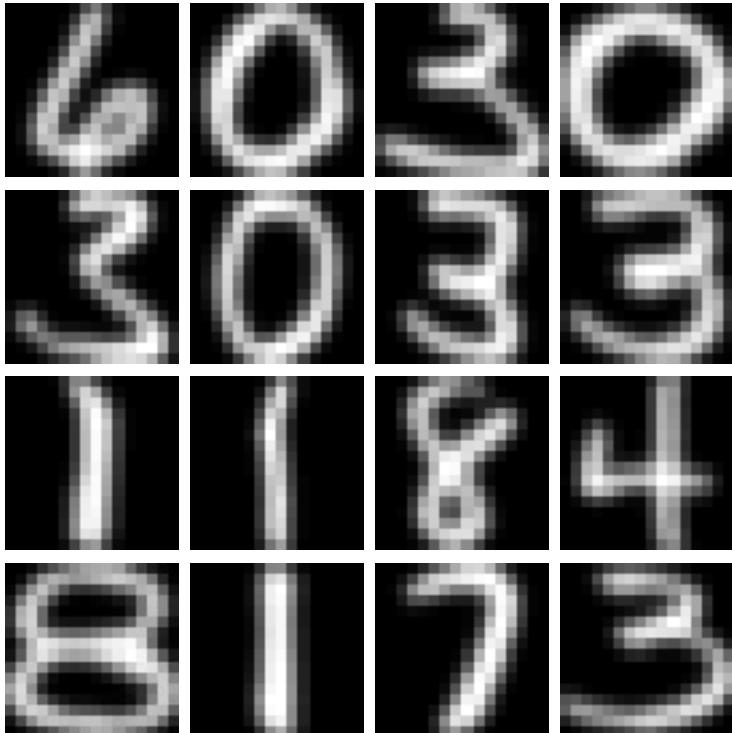
$$\frac{\partial \sigma_{\hat{\mathbf{w}}}^2}{\partial \mathbf{w}} = \frac{2}{\mathbf{w}^T \mathbf{w}} (\mathbf{C} \mathbf{w} - \sigma_{\hat{\mathbf{w}}}^2 \mathbf{w}) = 0 \implies$$

$$\mathbf{C} \mathbf{w} = \sigma_{\hat{\mathbf{w}}}^2 \mathbf{w}$$

PCA is the Eigen-value decomposition of
the data covariance matrix.

PCA - Example

About 9000 training examples



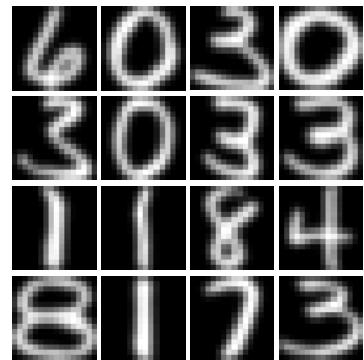
US Postal Service Digit Data

<http://www.gaussianprocess.org/gpml/data/>

Feature vectors

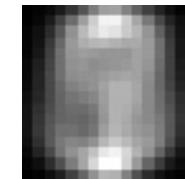
$$16 \times 16 \rightarrow \mathbf{x} = \begin{pmatrix} x_1 \\ \vdots \\ x_{256} \end{pmatrix}$$

PCA – Example, cont.



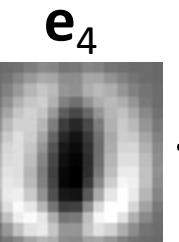
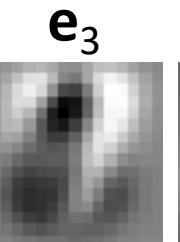
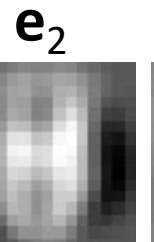
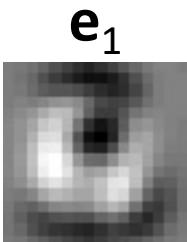
Mean digit

$$\bar{\mathbf{x}} =$$



$$\mathbf{C} = \frac{1}{N} \sum_{i=1}^N (\mathbf{x}_i - \bar{\mathbf{x}})(\mathbf{x}_i - \bar{\mathbf{x}})^T \quad 256 \times 256 \text{ matrix}$$

Eigendecomposition of \mathbf{C} :



...

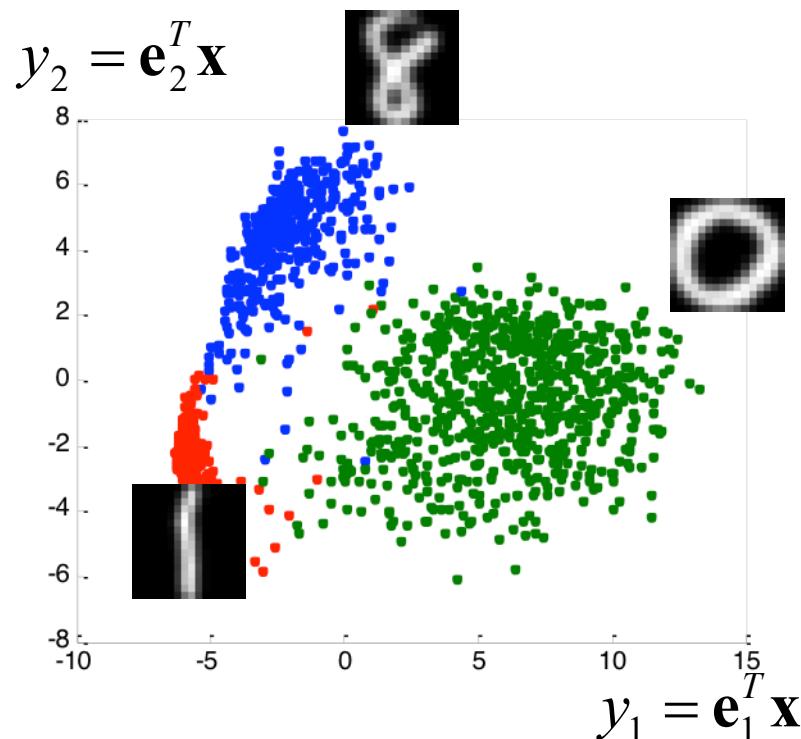
256 eigenvectors
and eigenvalues

50 100 150 200 250

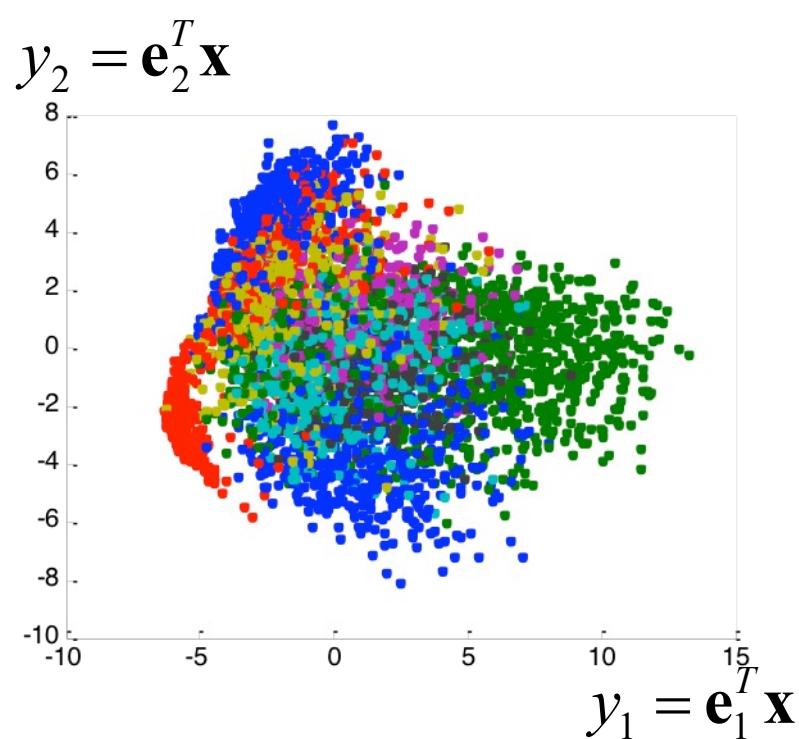
PCA – Example, cont.

From 256 to 2 feature dimensions!

3 classes



All 10 classes

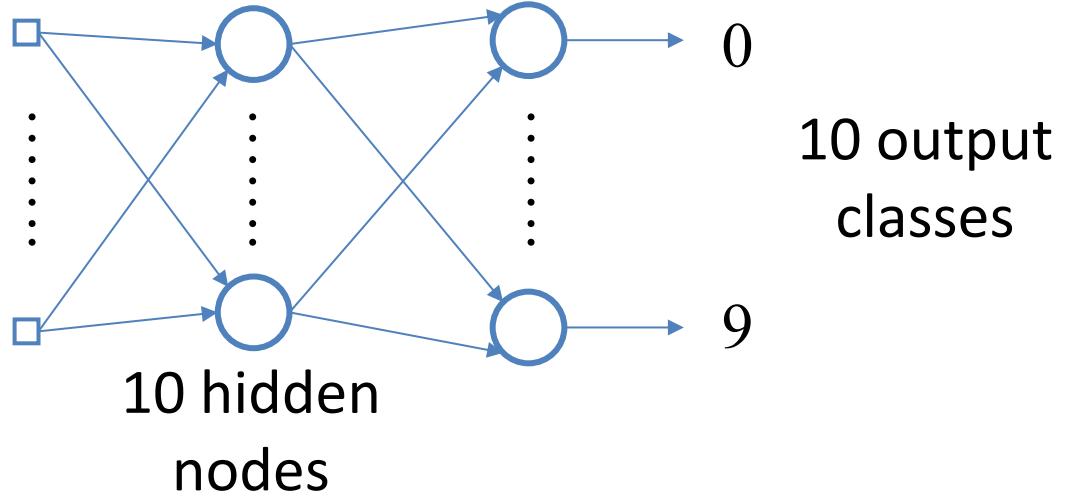


PCA – Example, cont.

Say we try to classify with a neural network with 10 hidden nodes



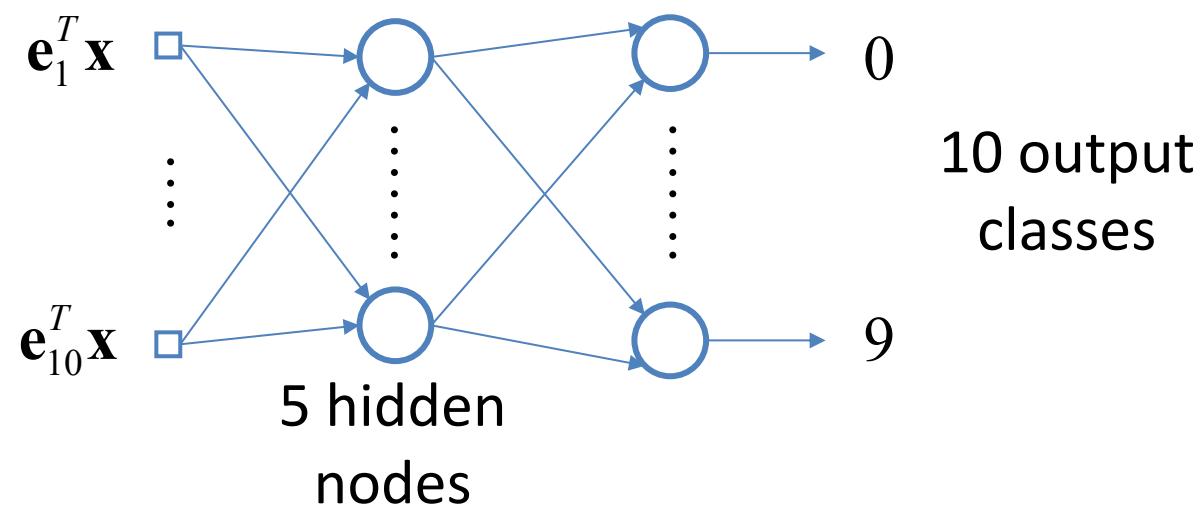
256 input
features



parameters in network: $257 \times 10 + 11 \times 10 = 2680$
(including bias weights)

PCA – Example, cont.

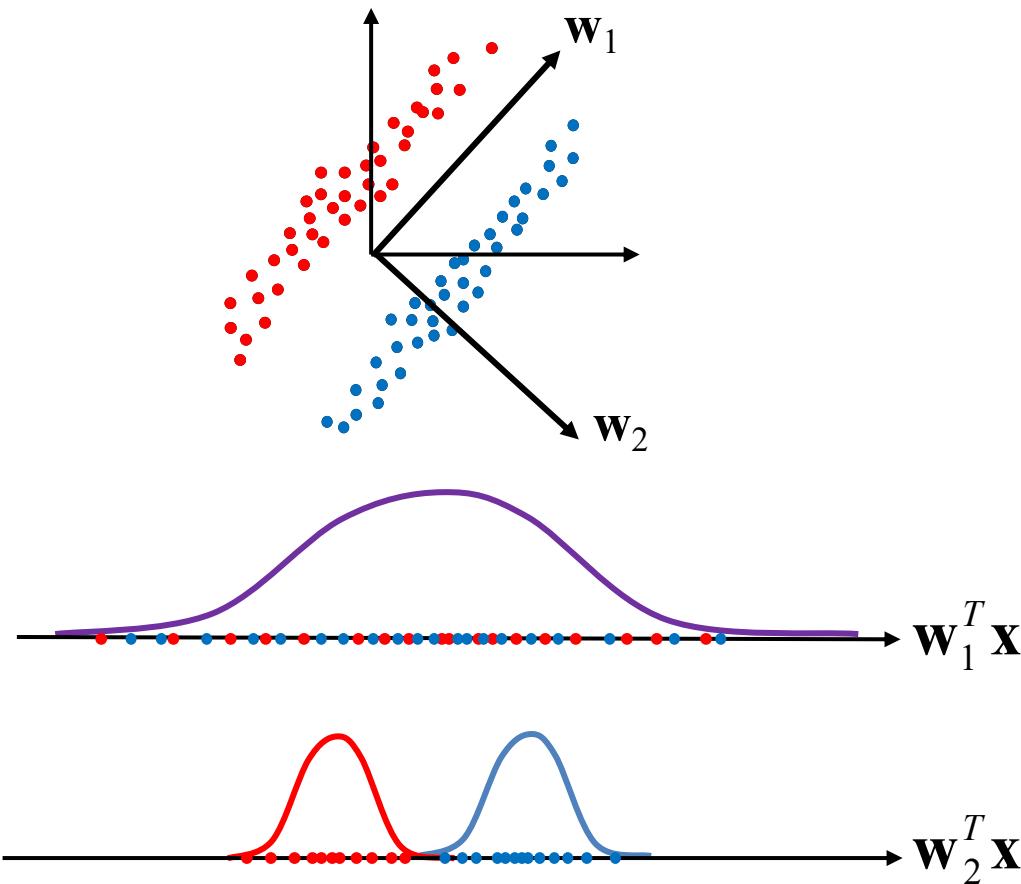
If we reduce the input dimensionality first, we may be able to do the classification with a smaller network, e.g., 10 principal components as input and 5 hidden nodes.



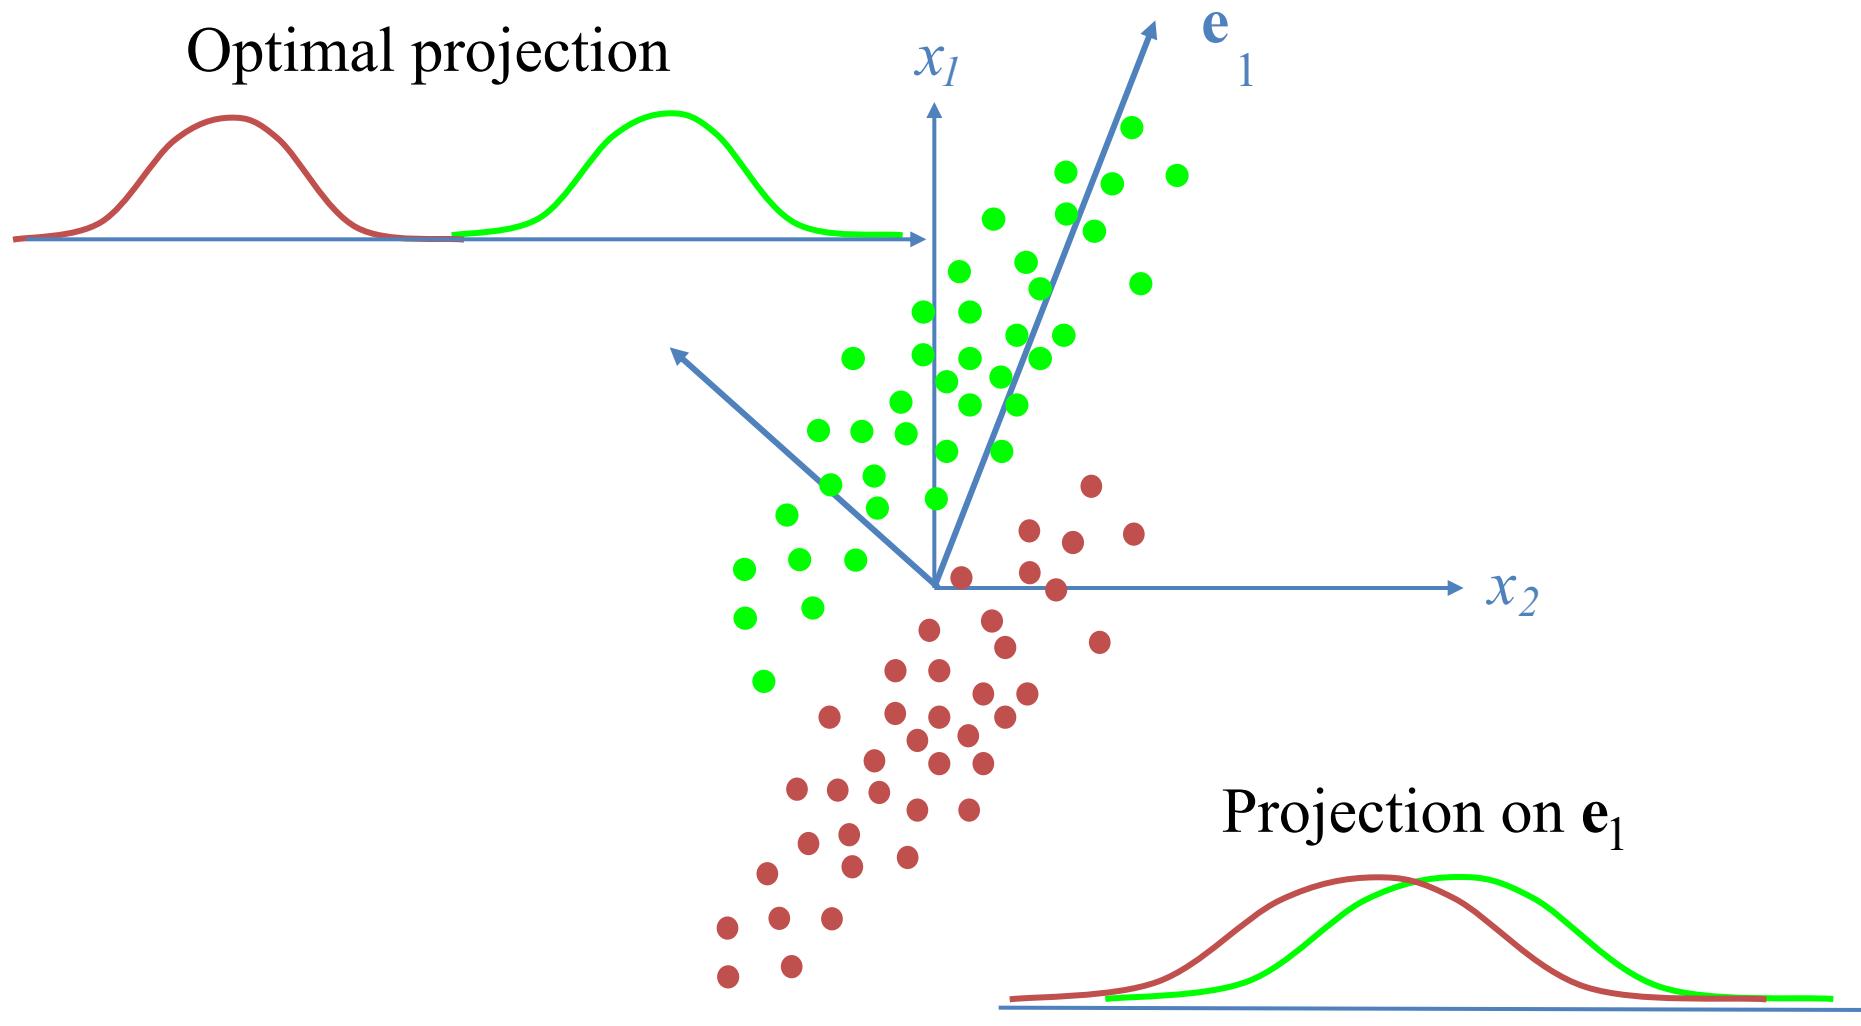
parameters in network: $11 \times 5 + 6 \times 10 = 115$
(including bias weights)

Limitations with PCA

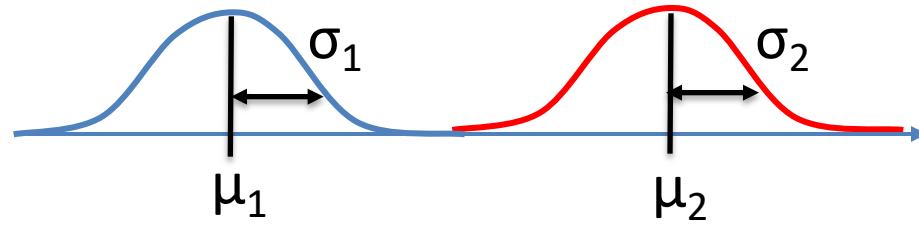
Variance is not always the most important goal!



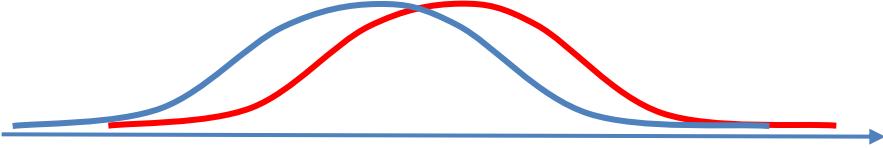
Optimal projection for separation of two clusters



Class separability



- Small variance
- Large distance



- Large variance
- Small distance

Goal: minimize variance and maximize distance.

Linear Discriminant Analysis (LDA)

a.k.a. Fishers Linear Discriminant (FLD)

- Minimize variance
- Maximize distance

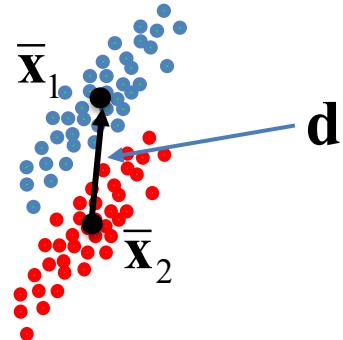
$$\text{Maximize: } \varepsilon(\mathbf{w}) = \frac{(\mu_1 - \mu_2)^2}{\sigma_1^2 + \sigma_2^2}$$

LDA – Loss function

$$\varepsilon = \frac{(\mu_1 - \mu_2)^2}{\sigma_1^2 + \sigma_2^2}$$

Distance:

$$\mu(\mathbf{w}) = \frac{1}{n} \sum_{i=1}^n \mathbf{w}^T \mathbf{x}_i = \mathbf{w}^T \left(\frac{1}{n} \sum_{i=1}^n \mathbf{x}_i \right) = \mathbf{w}^T \bar{\mathbf{x}}$$



$$(\mu_1(\mathbf{w}) - \mu_2(\mathbf{w}))^2 = (\mathbf{w}^T \bar{\mathbf{x}}_1 - \mathbf{w}^T \bar{\mathbf{x}}_2)^2 = (\mathbf{w}^T \mathbf{d})^2 = \mathbf{w}^T \underbrace{\mathbf{d} \mathbf{d}^T}_{\mathbf{M}} \mathbf{w} = \mathbf{w}^T \mathbf{M} \mathbf{w}$$

Variance:

$$\sigma^2(\mathbf{w}) = \frac{1}{n} \sum_{i=1}^n \left(\mathbf{w}^T (\mathbf{x}_i - \bar{\mathbf{x}}) \right)^2 = \dots = \mathbf{w}^T \underbrace{\left(\frac{1}{n} \sum_{i=1}^n (\mathbf{x}_i - \bar{\mathbf{x}})(\mathbf{x}_i - \bar{\mathbf{x}})^T \right)}_{\mathbf{C}} \mathbf{w} = \mathbf{w}^T \mathbf{C} \mathbf{w}$$

$$\sigma_1^2(\mathbf{w}) + \sigma_2^2(\mathbf{w}) = \mathbf{w}^T \mathbf{C}_1 \mathbf{w} + \mathbf{w}^T \mathbf{C}_2 \mathbf{w} = \mathbf{w}^T \underbrace{\mathbf{C}_{tot}}_{\mathbf{C}} \mathbf{w}$$

Exercise:
Complete all the steps!

Note the assumption $\mathbf{C}_1 = \mathbf{C}_2$!

LDA – Solution

$$\varepsilon(\mathbf{w}) = \frac{(\mu_1 - \mu_2)^2}{\sigma_1^2 + \sigma_2^2} = \frac{\mathbf{w}^T \mathbf{M} \mathbf{w}}{\mathbf{w}^T \mathbf{C}_{tot} \mathbf{w}}$$

Compare with PCA!

This form is called a Rayleigh quotient,
which is maximized by the largest eigenvector to the
generalized eigenvalue problem $\mathbf{C}_{tot}\mathbf{w} = \lambda \mathbf{M}\mathbf{w}$!

Simplification: $\mathbf{M}\mathbf{w} = \mathbf{d}\mathbf{d}^T \mathbf{w} = k\mathbf{d}$


Some scalar k

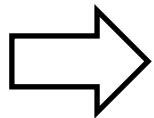
$$\mathbf{w} \sim \mathbf{C}_{tot}^{-1} \mathbf{d}$$

Scaling of \mathbf{w} not important!

LDA - Example

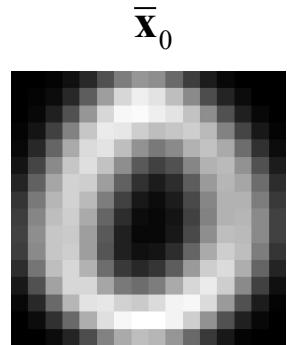


Feature vectors

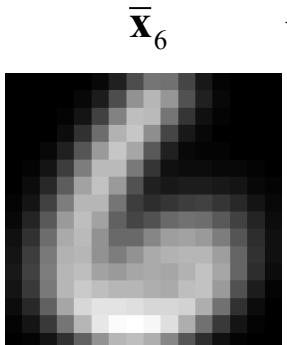


$$\mathbf{x}_0 = \begin{pmatrix} x_1 \\ \vdots \\ x_{256} \end{pmatrix}_{0\text{-digits}}$$

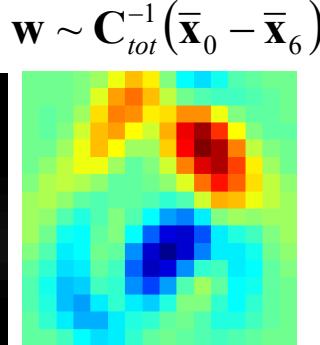
$$\mathbf{x}_6 = \begin{pmatrix} x_1 \\ \vdots \\ x_{256} \end{pmatrix}_{6\text{-digits}}$$



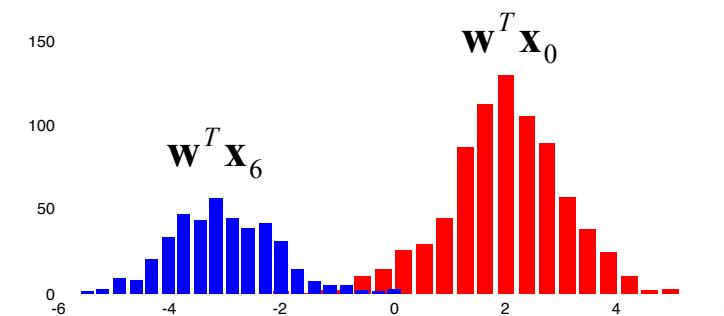
$\bar{\mathbf{x}}_0$



$\bar{\mathbf{x}}_6$

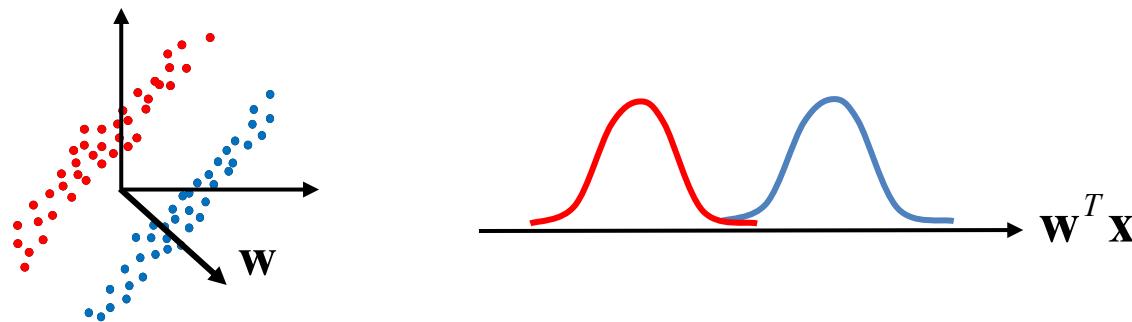


Reduction from 256 to 1 dimension



LDA - Summary

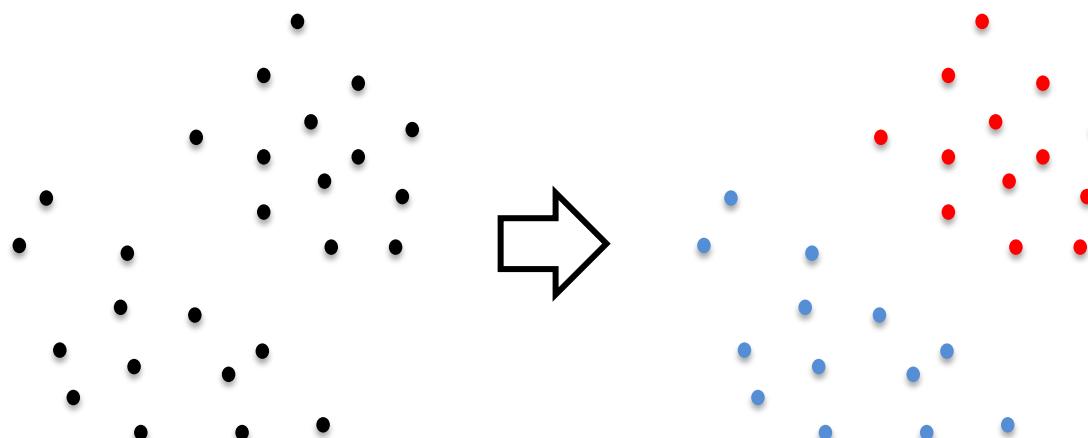
- Projection direction \mathbf{w} on which two classes are maximally separated



- Closed form solution, no parameters to set, no local optima.
- Components of \mathbf{w} give the importance of each feature in \mathbf{x} .
- Assumes identical distributions of the two classes!
- Note that this is actually supervised learning since labelled data is required!

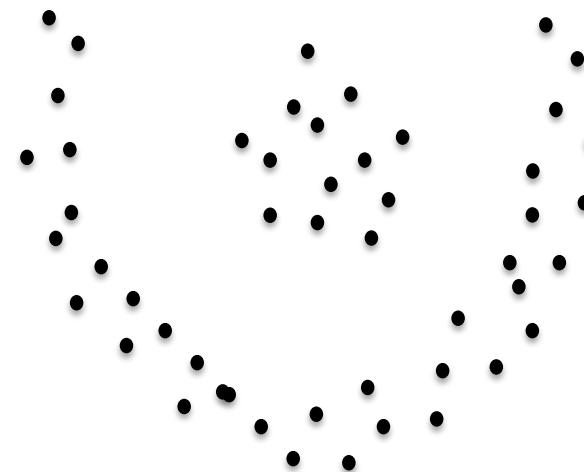
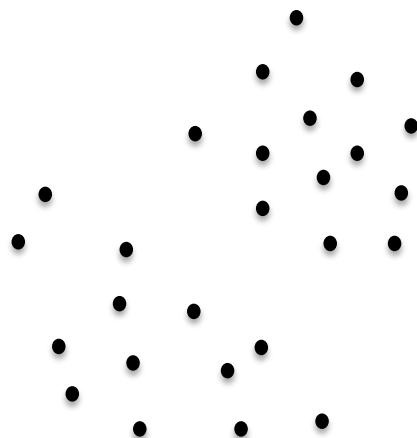
Categorization

- Categorization and grouping of objects based on similar properties is an important functionality in learning and knowledge representation.
- In the machine learning area, this is usually referred to as *clustering*.



What describes a cluster?

- Distances to other points?
- Connectivity?
- Different definitions lead to different algorithms.



k -Means algorithm

- Assume k clusters (user input).
- Represent each cluster with a mean prototype vector \mathbf{p}_j at the cluster center.
- A data point belongs to the cluster with the closest prototype vector (Euclidian distance).



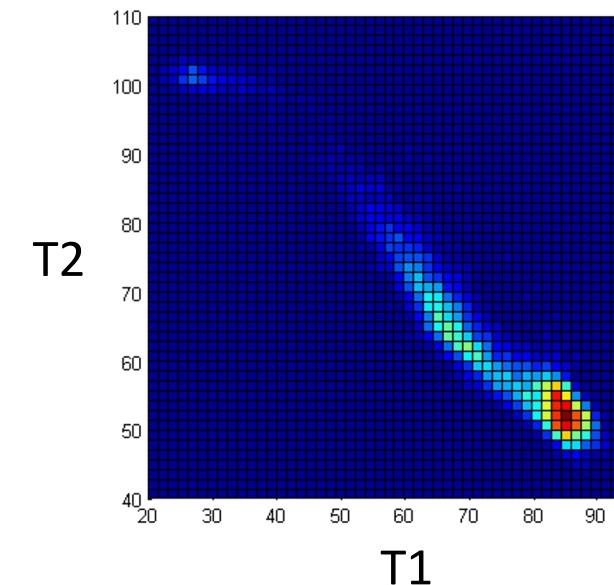
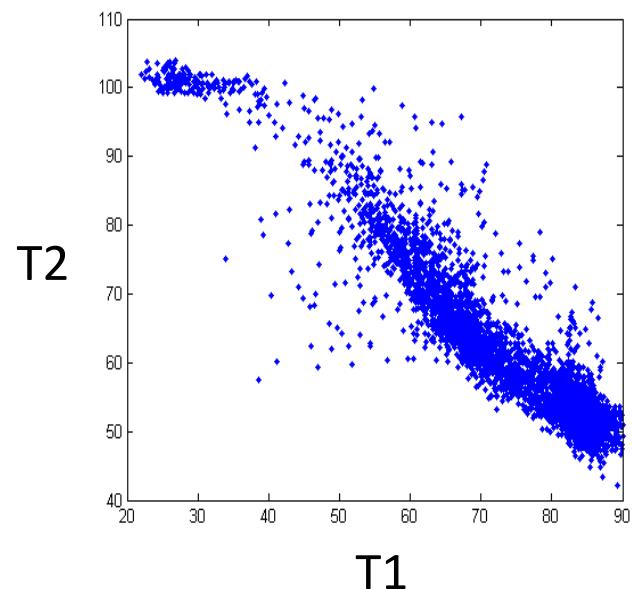
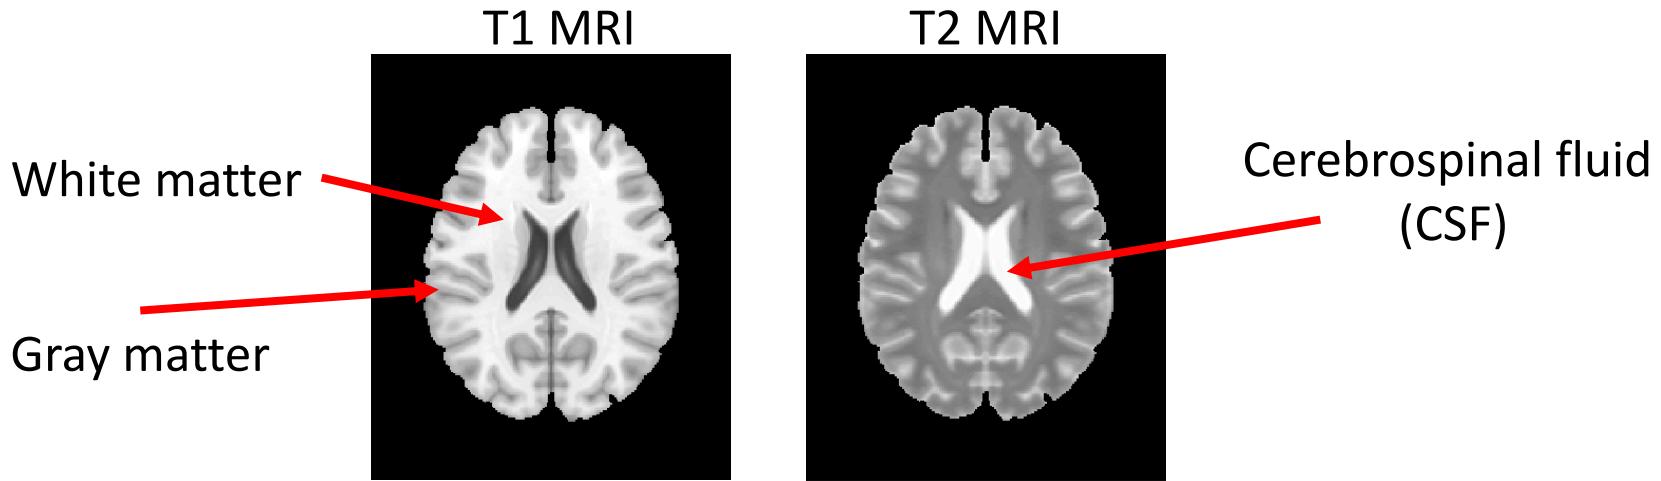
k -Means algorithm, cont.

1. Start with k random prototype vectors \mathbf{p}_j
2. Iterate:
 1. Assignment: Assign each data vector \mathbf{x}_i to the closest prototype vector \mathbf{p}_j . Denote the set of data vectors assigned to cluster \mathbf{p}_j by S_j .
 2. Update all prototype vectors to the mean of the clusters:

$$\mathbf{p}_j = \frac{1}{|S_j|} \sum_{k \in S_j} \mathbf{x}_k$$

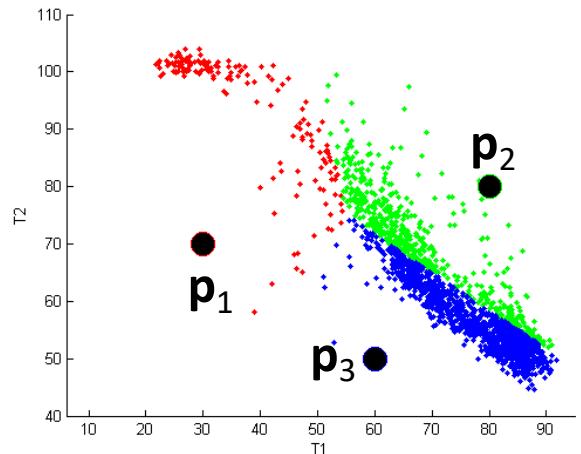
Number of elements in S

k -Means - Example

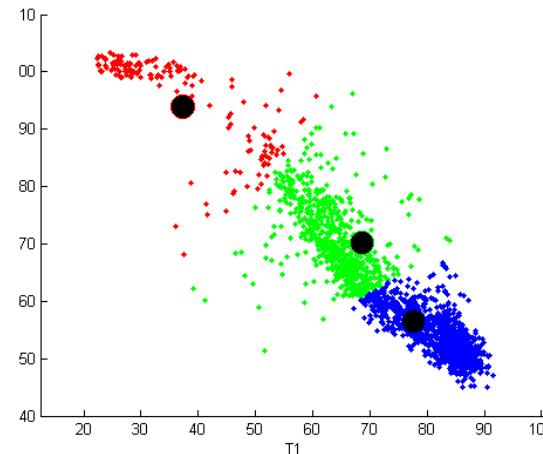


k -Means - Example

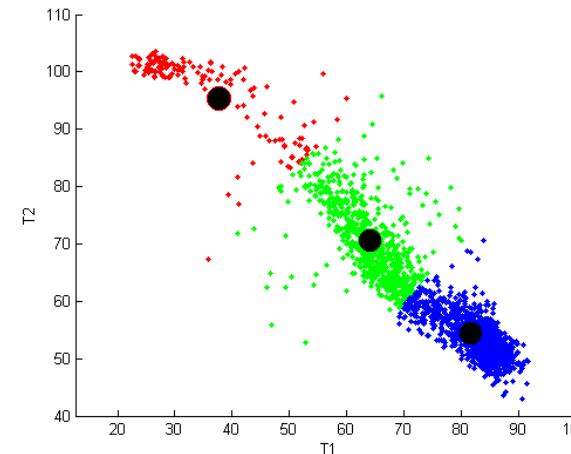
Random initialization



1st iteration



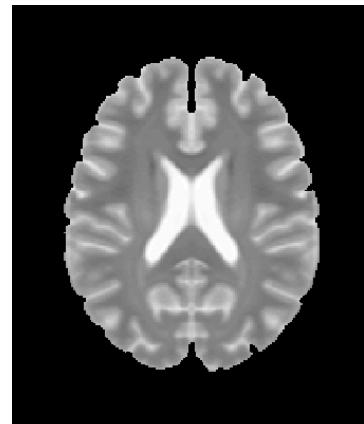
2nd iteration



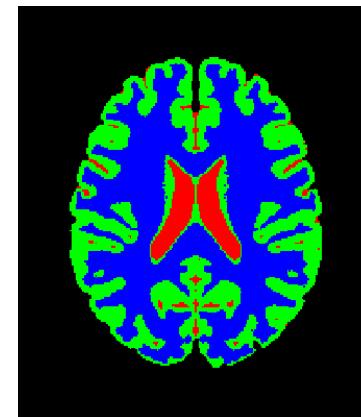
T1 MRI



T2 MRI



k -means result



k -Means - Discussion

- Must specify k
- Tries to minimize the loss function

$$\varepsilon(\mathbf{p}_1, \dots, \mathbf{p}_k, S_1, \dots, S_k) = \sum_{i=1}^k \sum_{\mathbf{x}_j \in S_i} \|\mathbf{x}_j - \mathbf{p}_i\|^2$$

- Note that this function is not differentiable as the S_i :s are discrete sets, i.e., we cannot do gradient descent.

Expectation Maximization (EM) – Intro

$$\varepsilon(\mathbf{p}_1, \dots, \mathbf{p}_k, \underbrace{S_1, \dots, S_k}_{\text{Unknown class labels}}) = \sum_{i=1}^k \sum_{\mathbf{x}_j \in S_i} \|\mathbf{x}_j - \mathbf{p}_i\|^2$$

Assume we know S_1, \dots, S_k !

$$\frac{\partial \varepsilon}{\partial \mathbf{p}_i} = \frac{\partial}{\partial \mathbf{p}_i} \left(\sum_{\mathbf{x}_j \in S_i} \|\mathbf{x}_j - \mathbf{p}_i\|^2 \right) = -2 \sum_{\mathbf{x}_j \in S_i} (\mathbf{x}_j - \mathbf{p}_i) = 2|S_i| \mathbf{p}_i - 2 \sum_{\mathbf{x}_j \in S_i} \mathbf{x}_j$$

$$\frac{\partial \varepsilon}{\partial \mathbf{p}_i} = 0 \quad \xrightarrow{\text{Mean vector!}} \quad \mathbf{p}_i = \frac{1}{|S_i|} \sum_{\mathbf{x}_j \in S_i} \mathbf{x}_j$$

Expectation Maximization (EM) – Intro

$$\varepsilon \left(\underbrace{\mathbf{p}_1, \dots, \mathbf{p}_k}_{\text{Continuous parameters}}, S_1, \dots, S_k \right) = \sum_{i=1}^k \sum_{\mathbf{x}_j \in S_i} \|\mathbf{x}_j - \mathbf{p}_i\|^2$$

Now, assume we know $\mathbf{p}_1, \dots, \mathbf{p}_k$!

Each \mathbf{x}_j independently contributes a distance $\|\mathbf{x}_j - \mathbf{p}_i\|$ to ε

Obvious that ε is minimized if each \mathbf{x}_j is assigned to the set S associated with the closest \mathbf{p} !

k-Means algorithm!!

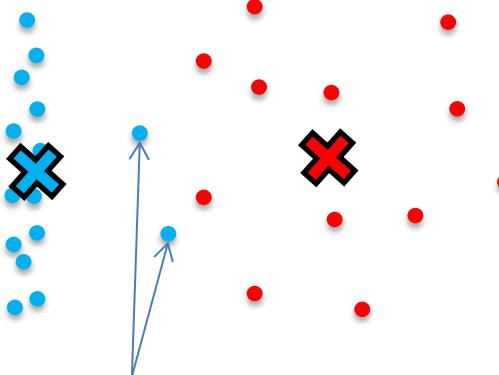
Expectation Maximization

- The k -means algorithm is a special case of a general optimization approach called *Expectation Maximization* (EM).
- EM can be used when we want to estimate model parameters (the prototypes \mathbf{p}_i), but for each data sample \mathbf{x}_j there is a hidden/missing discrete parameter (the class labels S_j).
- EM iterates between optimizing the hidden parameters and the model parameters.

Mixture of Gaussians (MoG) clustering

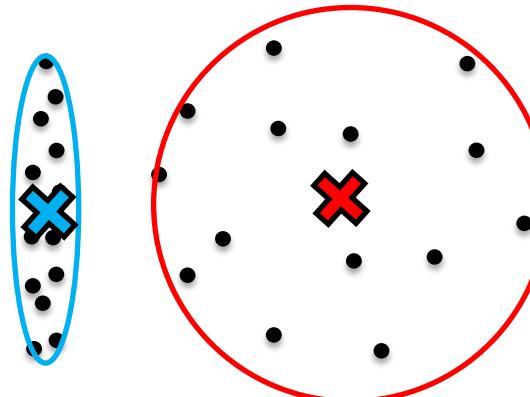
Another application of EM!

k-means



May not be correctly clustered!

Idea: Represent each cluster using a Gaussian distribution



Problem: Each data sample \mathbf{x}_j belongs to one of k Gaussian distributions $N(\mathbf{p}_i, \mathbf{C}_i)$, $i=1..k$.

Find the sets S_i of samples that belong to distribution $N(\mathbf{p}_i, \mathbf{C}_i)$ and the mean \mathbf{p}_i and covariance matrix \mathbf{C}_i of that distribution.

Let's use EM!

Assume we know the hidden parameters S_1, \dots, S_k !

That is, we know the samples for each set, e.g., $S_1 = \{\mathbf{x}_1, \mathbf{x}_7, \mathbf{x}_{12}, \mathbf{x}_{13}\}$

We can then estimate the mean \mathbf{p}_i and covariance \mathbf{C}_i using standard estimation for the Gaussian:

$$\mathbf{p}_i = \frac{1}{|S_i|} \sum_{k \in S_i} \mathbf{x}_k$$

$$\mathbf{C}_i = \frac{1}{|S_i|} \sum_{k \in S_i} (\mathbf{x}_k - \mathbf{p}_i)(\mathbf{x}_k - \mathbf{p}_i)^T$$

Let's use EM, cont!

Assume now that we know the Gaussian distribution parameters, i.e., we know all distributions

$$f(\mathbf{x}; \mathbf{p}_i, \mathbf{C}_i) = \frac{1}{(2\pi)^{d/2} |\mathbf{C}_i|^{1/2}} e^{-\frac{1}{2}(\mathbf{x}-\mathbf{p}_i)^T \mathbf{C}_i^{-1} (\mathbf{x}-\mathbf{p}_i)}$$

Let the hidden parameter S_i be the set of all data samples for which $f(\mathbf{x}; \mathbf{p}_i, \mathbf{C}_i)$ is larger than for all other distributions. That is, each data sample is assigned to the Gaussian distribution **to which it most likely belongs!**

Mixture of Gaussians – Using EM

1. Start with k random Gaussians $(\mathbf{p}_j, \mathbf{C}_j)$
2. Iterate:
 1. Assign each sample \mathbf{x}_i to the most likely Gaussian

$$\max f(\mathbf{x}_i; \mathbf{p}_j, \mathbf{C}_j) = \frac{1}{(2\pi)^{d/2} |\mathbf{C}_j|^{1/2}} e^{-\frac{1}{2} (\mathbf{x}_i - \mathbf{p}_j)^T \mathbf{C}_j^{-1} (\mathbf{x}_i - \mathbf{p}_j)}$$

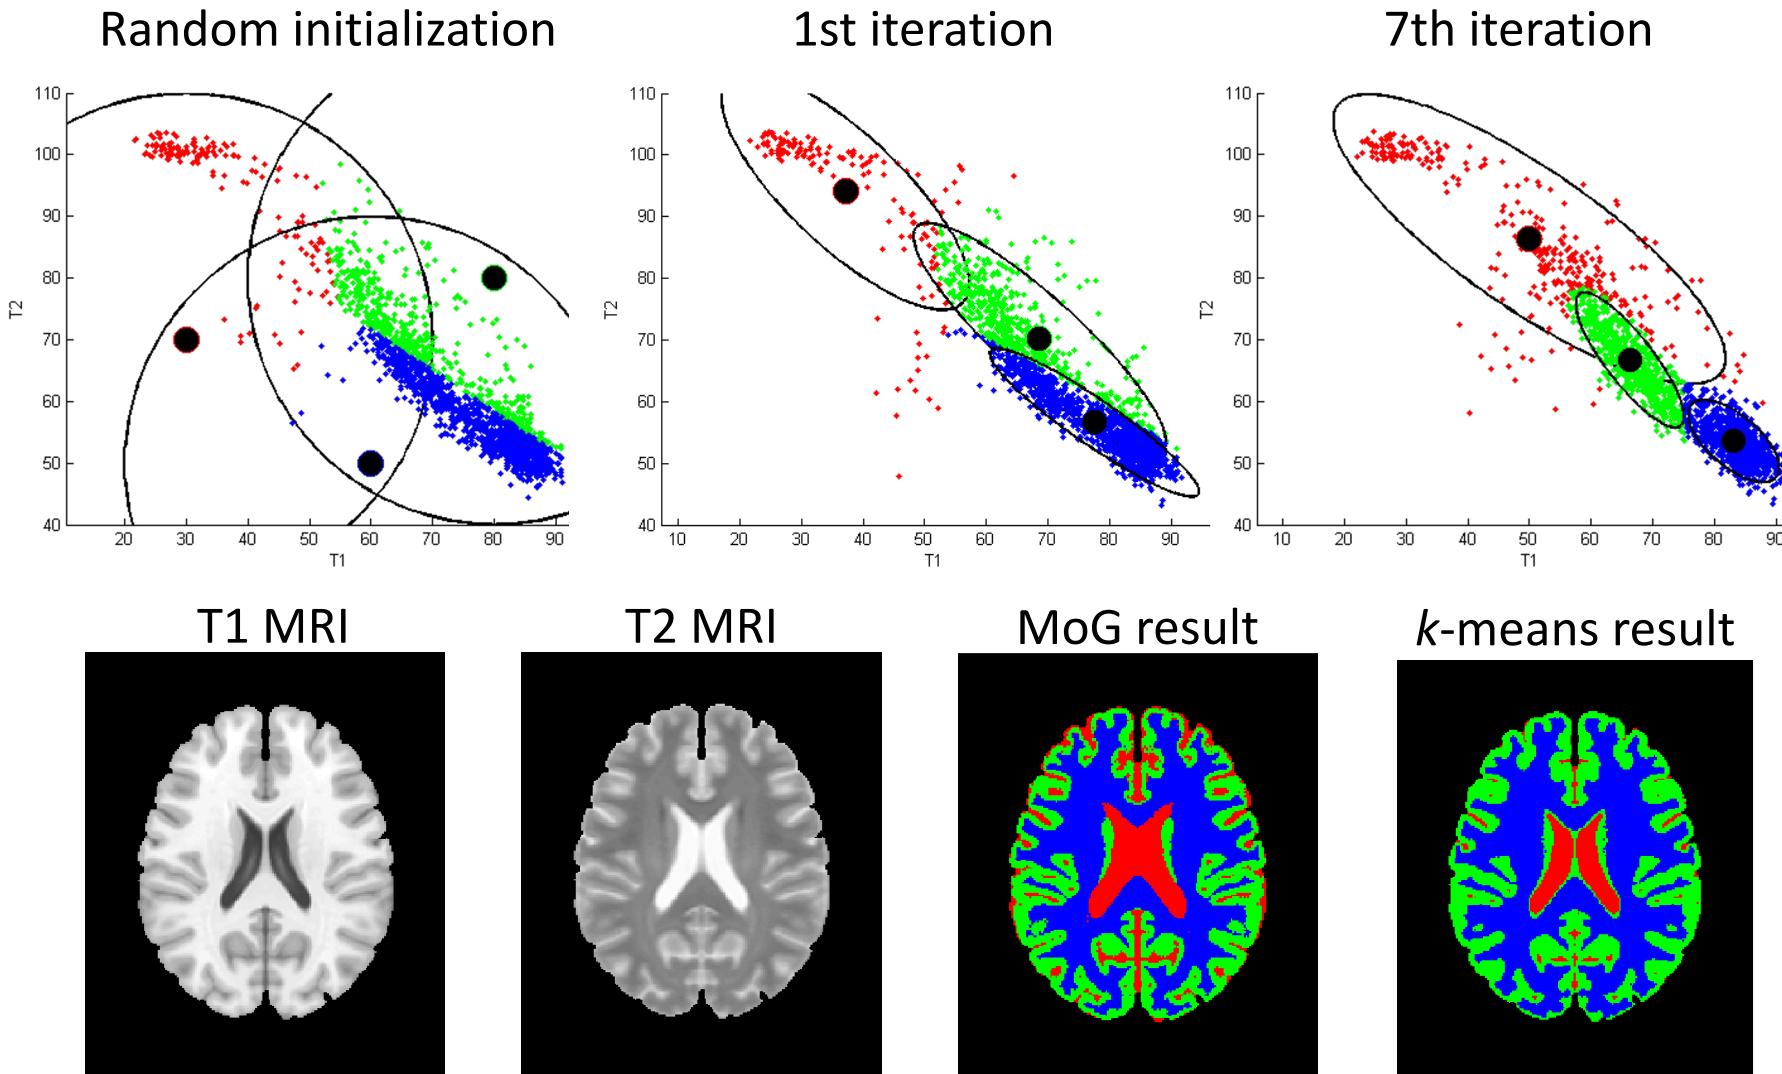
Denote the set of samples assigned to Gaussian j by S_j .

2. Update all Gaussian's means and covariances:

$$\mathbf{p}_j = \frac{1}{|S_j|} \sum_{k \in S_j} \mathbf{x}_k$$

$$\mathbf{C}_j = \frac{1}{|S_j|} \sum_{k \in S_j} (\mathbf{x}_k - \mathbf{p}_j)(\mathbf{x}_k - \mathbf{p}_j)^T$$

Mixture of Gaussians - Example



Summary of k -Means and MoG clustering

- Must choose the number of clusters k .
- Different initializations may give different results.
- May converge to degenerate solutions, e.g., empty clusters.
- MoG allows for elliptic cluster shapes with different sizes.

Auto Encoders

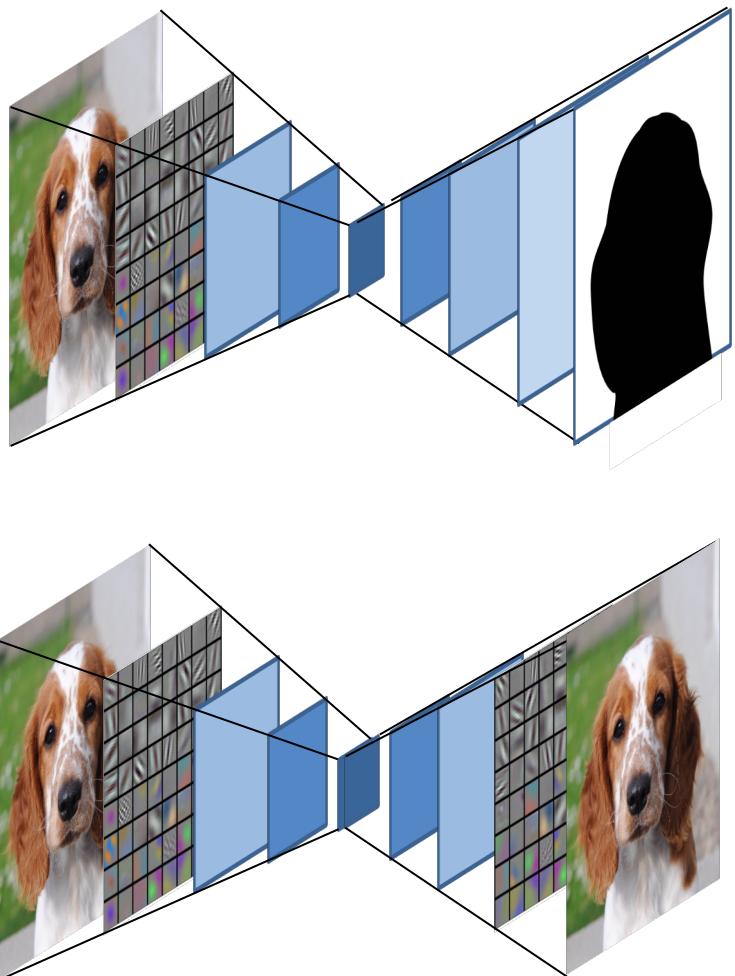


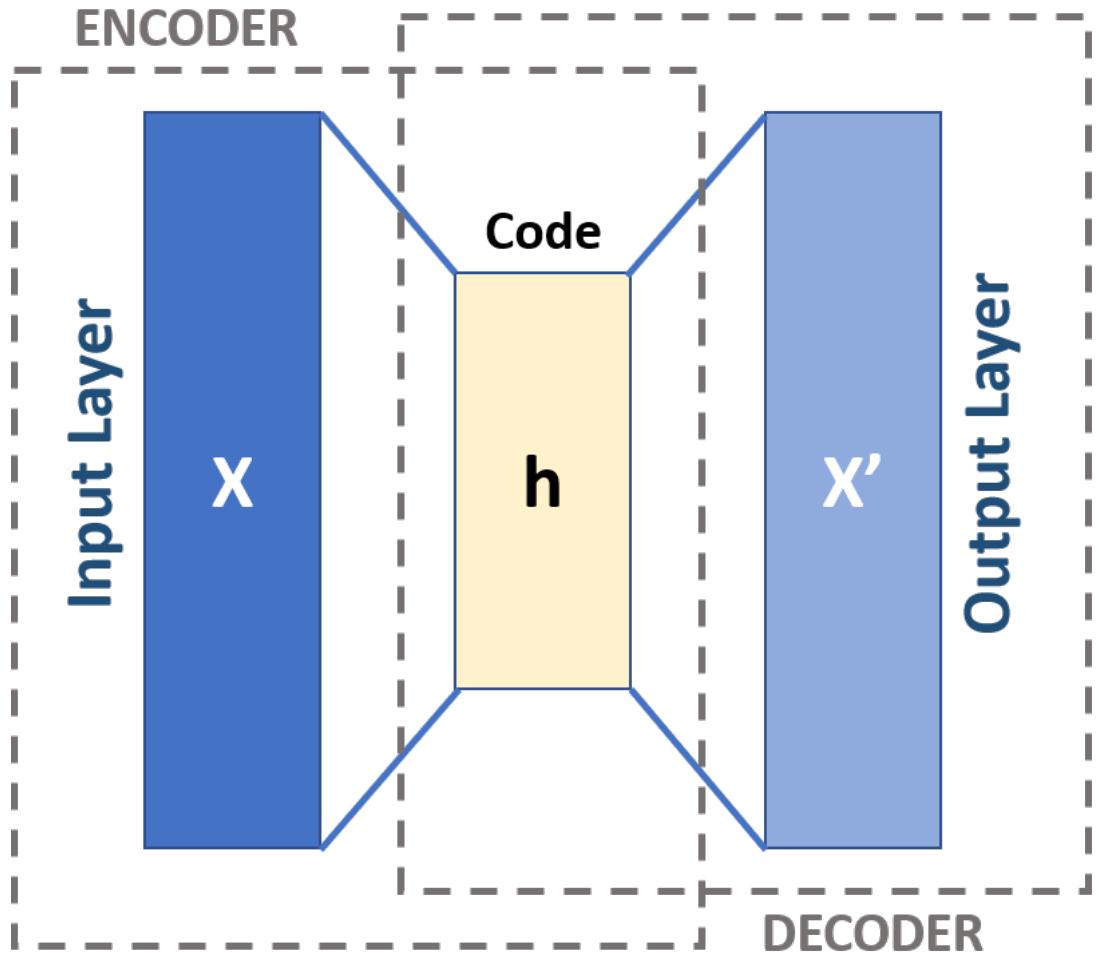
Image segmentation

- Training data $\{x_i, y_i\}$
- Binary output (foreground / background)
- Classification of each pixel

Auto encoder

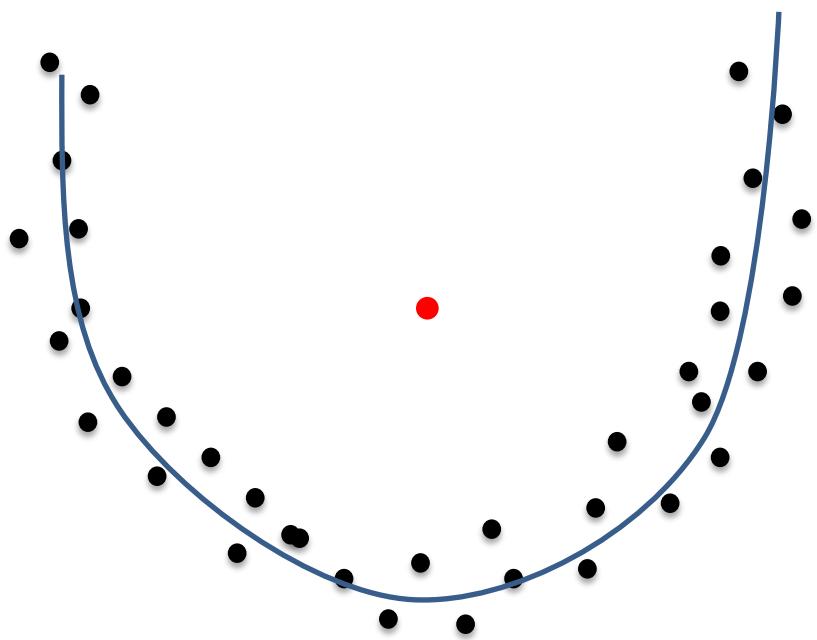
- Training data $\{x_i, x_i\}$
- Does not require labelled training data – Unsupervised learning!
- The smallest layer must be able to represent all training data
- The bottle neck forces an efficient coding of the original data (images)

Auto Encoders



- The hidden representation h is a compact encoding of the distribution of training data x
- If the dimensionality of h is chosen properly, h can represent all training data well, but not data points that deviate from the distribution of x

Data manifold



- The data lies in a low-dimensional manifold in a high-dimensional feature space.
- An auto encoder learns to represent the manifold
- Data points that are outside the manifold cannot be accurately represented

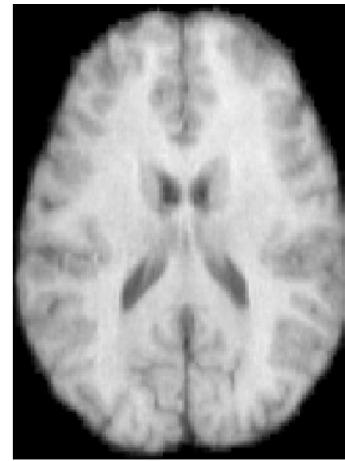
Outlier detection



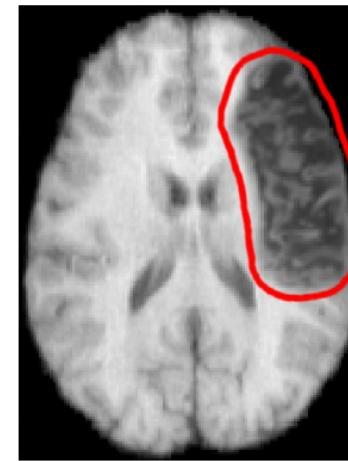
- Data points that deviate from the distribution of normal samples (the learned manifold)
- These data points can indicate anomalies such as production errors or pathologies

Application in Medical Image Analysis

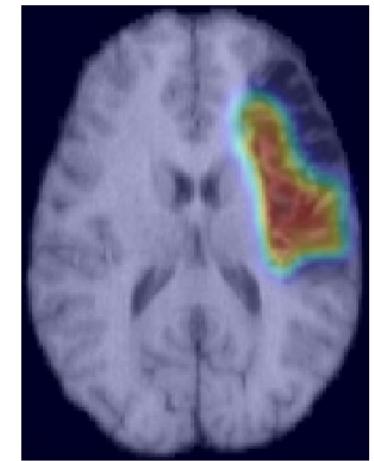
- Train on “normal images”
- Compare input and output
- Use the residual as indication of abnormality



Normal
brain



Brain with
lesion



Residual
overlay